



# Balanced Search Trees



Computer Science 112  
Boston University


Christine Papadakis-Kanaris

# Recall: Path, Depth, Level, and Height




- There is exactly one *path* (one sequence of edges) connecting each node to the root.
- *depth* of a node = # of edges on the path from it to the root
- Nodes with the same depth form a *level* of the tree.
- The *height* of a tree is the ***maximum depth of its nodes***.
  - example: the tree above has a height of 2

# Recall: A Tree is a Recursive Data Structure




- **Each node in the tree is the **root** of a smaller tree!**
  - refer to such trees as *subtrees* to distinguish them from the tree as a whole
  - example: node 2 is the root of the subtree circled above
  - example: node 6 is the root of a subtree with only one node

# Recall: A Tree is a Recursive Data Structure




- Height 3

# Recall: A Tree is a Recursive Data Structure




- Height 2

# Recall: A Tree is a Recursive Data Structure



- Height 0


# Recall: A Tree is a Recursive Data Structure



- What is the height of empty nodes? -1

# Defining Balanced Trees


- A tree is *balanced* if, for *each* of its nodes, the node's subtrees have the same height or have heights that differ by 1.
  - example:
    - 26: both subtrees have a height of 1






# Defining Balanced Trees

- A tree is *balanced* if, for *each* of its nodes, the node's subtrees have the same height or have heights that differ by 1.
  - example:
    - 26: both subtrees have a height of 1
    - 12: left subtree has height 0




# Defininf Balanced Trees

- A tree is *balanced* if, for *each* of its nodes, the node's subtrees have the same height or have heights that differ by 1.
  - example:
    - 26: both subtrees have a height of 1
    - 12: left subtree has height 0  
right subtree is empty (height = -1)



# Defining Balanced Trees

- A tree is *balanced* if, for *each* of its nodes, the node's subtrees have the same height or have heights that differ by 1.
  - example:
    - 26: both subtrees have a height of 1
    - 12: left subtree has height 0  
right subtree is empty (height = -1)
    - 32: both subtrees have a height of 0





# Defining Balanced Trees

- A tree is *balanced* if, for *each* of its nodes, the node's subtrees have the same height or have heights that differ by 1.

- example:


- 26: both subtrees have a height of 1
- 12: left subtree has height 0  
right subtree is empty (height = -1)
- 32: both subtrees have a height of 0
- all leaf nodes: both subtrees are empty



# Efficiency of Balanced Trees

- A tree is *balanced* if, for *each* of its nodes, the node's subtrees have the same height or have heights that differ by 1.

- example:
  - 26: both subtrees have a height of 1
  - 12: left subtree has height 0  
right subtree is empty (height = -1)
  - 32: both subtrees have a height of 0
  - all leaf nodes: both subtrees are empty




- For a balanced tree with  $n$  nodes, height =  $O(\log n)$ 
  - each time that you follow an edge down the longest path, you cut the problem size roughly in half!
- Therefore, for a *balanced binary search tree*, the worst case for search / insert / delete is  $O(h) = O(\log n)$ 
  - the "best" worst-case time complexity

# 2-3 Trees


Insert in a 2-node :



Insert in a 3-node (2 node parent) :




Insert in a 3-node (3 node parent) :



## 2-3 Trees

- A 2-3 tree is a balanced tree in which:
  - *all* nodes have equal-height subtrees (perfect balance)
  - each node is either
    - a **2-node**, which contains **one** data item and *exactly* **0 or 2** children


- Example:



## 2-3 Trees

- A 2-3 tree is a balanced tree in which:
  - *all* nodes have equal-height subtrees (perfect balance)
  - each node is either
    - a **2-node**, which contains one data item and exactly 0 or 2 children
    - a **3-node**, which contains **two** data items and *exactly* 0 or 3 children

- Example:






# 2-3 Trees

- A 2-3 tree is a balanced tree in which:
  - *all* nodes have equal-height subtrees
  - each node is either
    - a **2-node**, which contains one key
    - a **3-node**, which contains two keys and two children


Assume the data in the node are the keys then...

- Example:




# 2-3 Trees


- A 2-3 tree is a balanced tree in which:
  - *all* nodes have equal-height subtrees (perfect balance)
  - each node is either
    - a **2-node**, which contains one data item and exactly 0 or 2 children
    - a **3-node**, which contains two data items and 0 or 3 children
  - the keys form a **search** tree
- Example:



2-node:




3-node:




# 2-3 Trees

- A 2-3 tree is a balanced tree in which:
  - *all* nodes have equal-height subtrees (perfect balance)
  - each node is either
    - a **2-node**, which contains one data item and 0 or 2 children
    - a **3-node**, which contains two data items and 0 or 3 children
  - the keys form a search tree


- Example:



2-node:




3-node:




# 2-3 Trees

- A 2-3 tree is a balanced tree in which:
  - *all* nodes have equal-height subtrees (perfect balance)
  - each node is either
    - a **2-node**, which contains one data item and 0 or 2 children
    - a **3-node**, which contains two data items and 0 or 3 children
  - the keys form a search tree


- Example:



2-node:



3-node:




# 2-3 Trees


- A 2-3 tree is a balanced tree in which:
  - *all* nodes have equal-height subtrees
  - each node is either
    - a **2-node**, which contains one data item
    - a **3-node**, which contains two data items
  - the keys form a search tree

Consider how the  
**two – three** constraint  
forces perfect balance!


- Example:



2-node:




3-node:



## 2-3 Trees


Add an A into our 2-3 tree:



# 2-3 Trees

Add an A into our 2-3 tree:

Add a B into our 2-3 tree:




# 2-3 Trees

Add an A into our 2-3 tree:

Add a B into our 2-3 tree:

Add a C into our 2-3 tree?



Can only have ***one***  
or ***two*** data items!




# 2-3 Trees

Add an A into our 2-3 tree:

Add a B into our 2-3 tree:

Add a C into our 2-3 tree?




# 2-3 Trees

Add an A into our 2-3 tree:

Add a B into our 2-3 tree:

Add a C into our 2-3 tree?




**Split the node!**

# 2-3 Trees

Add an A into our 2-3 tree:

Add a B into our 2-3 tree:

Add a C into our 2-3 tree?




# 2-3 Trees

Add an A into our 2-3 tree:

Add a B into our 2-3 tree:

Add a C into our 2-3 tree?



## 2-3 Trees


Add an A into

Perfect  
Balance!

2-node  
two children

Add a B into our 2-3 tree:

Add a C into our 2-3 tree?



two 2-nodes  
No children

## 2-3 Trees

Add an A into

Need to maintain  
the property of the  
search tree!

2-node  
two children

< B

B

>= B




Add a B into our 2-3 tree:

A

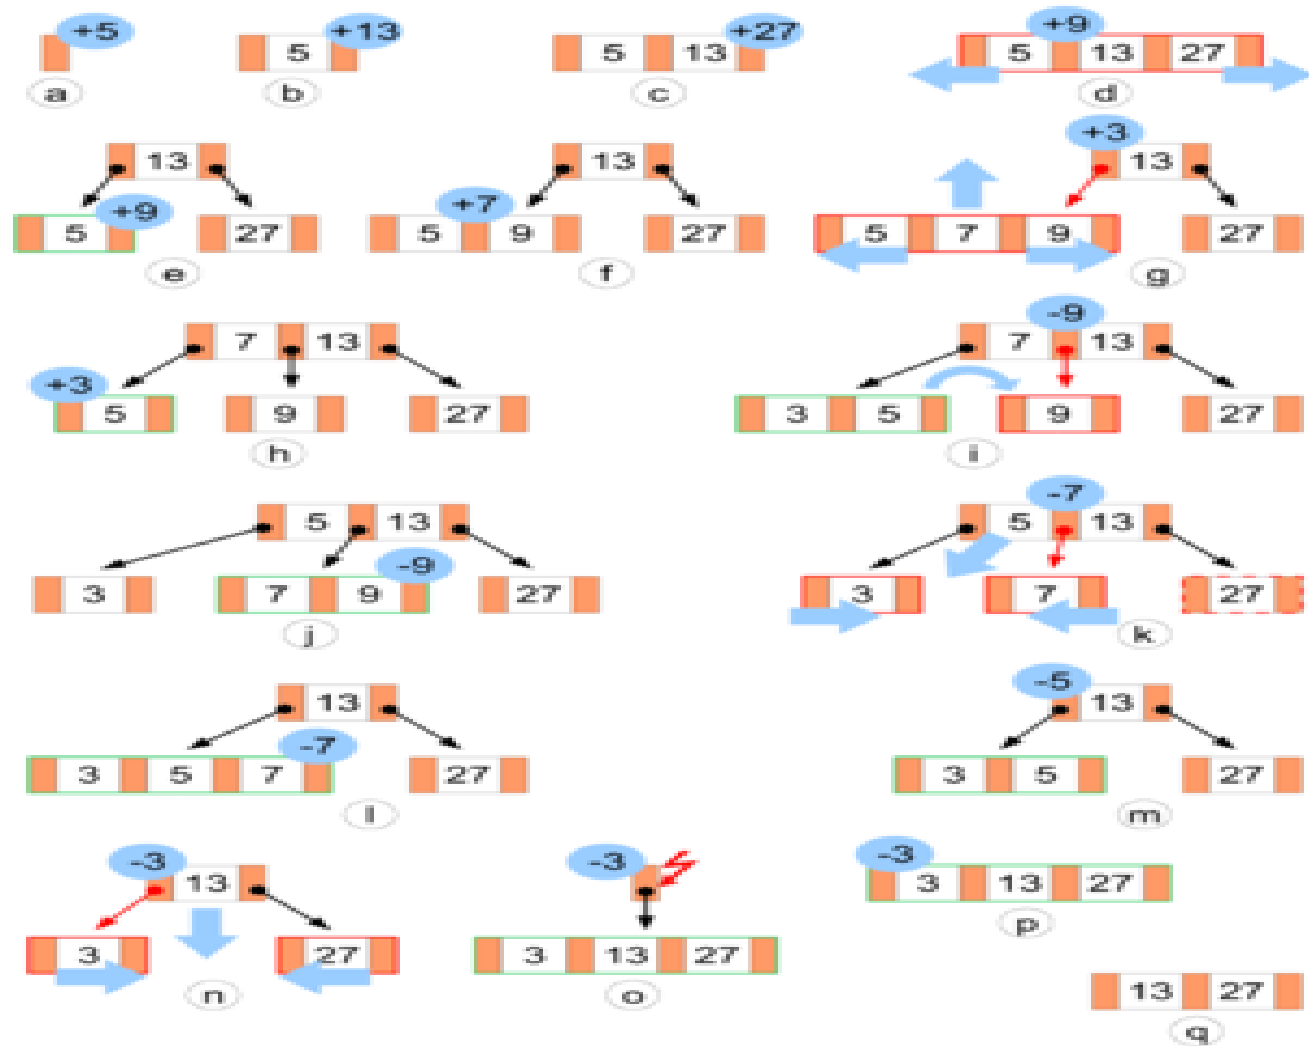
C

Add a C into our 2-3 tree?

two 2-nodes  
No children




# Operations on 2-3 trees




# Search in 2-3 Trees

- Algorithm for searching for an item with a key  $k$ :
  - if  $k ==$  one of the root node's keys, you're done
  - else if  $k <$  the root node's first key
    - search the left subtree
  - else if the root is a 3-node and  $k <$  its second key
    - search the middle subtree
  - else
    - search the right subtree



- Example: search for 87






# Insertion in 2-3 Trees

- Algorithm for inserting an item with a key  $k$ :  
search for  $k$ , but don't stop until you hit a leaf node  
let  $L$  be the leaf node at the end of the search

if  $L$  is a 2-node

add  $k$  to  $L$ , making it a 3-node

*example : add 14*




# Insertion in 2-3 Trees

- Algorithm for inserting an item with a key  $k$ :  
search for  $k$ , but don't stop until you hit a leaf node  
let  $L$  be the leaf node at the end of the search

if  $L$  is a 2-node

add  $k$  to  $L$ , making it a 3-node

*example : add 14*



How are we able to turn  
this 2-node into a .....


# Insertion in 2-3 Trees

- Algorithm for inserting an item with a key  $k$ :  
search for  $k$ , but don't stop until you hit a leaf node  
let  $L$  be the leaf node at the end of the search

if  $L$  is a 2-node

add  $k$  to  $L$ , making it a 3-node

*example* : add 14



How are we able to turn  
this 2-node into a  
3-node?


# Insertion in 2-3 Trees

- Algorithm for inserting an item with a key  $k$ :  
search for  $k$ , but don't stop until you hit a leaf node  
let  $L$  be the leaf node at the end of the search

if  $L$  is a 2-node

add  $k$  to  $L$ , making it a 3-node


*example* : add 14



Both a 2-node and a 3-node can have **no** children!

# Insertion in 2-3 Trees


- Algorithm for inserting an item with a key  $k$ :
  - search for  $k$ , but don't stop until you hit a leaf node
  - let  $L$  be the leaf node at the end of the search
  - if  $L$  is a 2-node
    - add  $k$  to  $L$ , making it a 3-node



else if  $L$  is a 3-node

split  $L$  into two 2-nodes containing the items with the smallest and largest of:  $k$ ,  $L$ 's 1<sup>st</sup> key,  $L$ 's 2<sup>nd</sup> key  
the middle item is “sent up” and inserted in  $L$ 's parent

*example: add 52*





# Insertion in 2-3 Trees

- Algorithm for inserting an item with a key  $k$ :
  - search for  $k$ , but don't stop until you hit a leaf node
  - let  $L$  be the leaf node at the end of the search
  - if  $L$  is a 2-node
    - add  $k$  to  $L$ , making it a 3-node

else if  $L$  is a 3-node

- split  $L$  into two 2-nodes
- the smallest and largest keys go to the new left and right children
- the middle item is inserted in  $L$ 's parent

*example: add 52*



# Insertion in 2-3 Trees

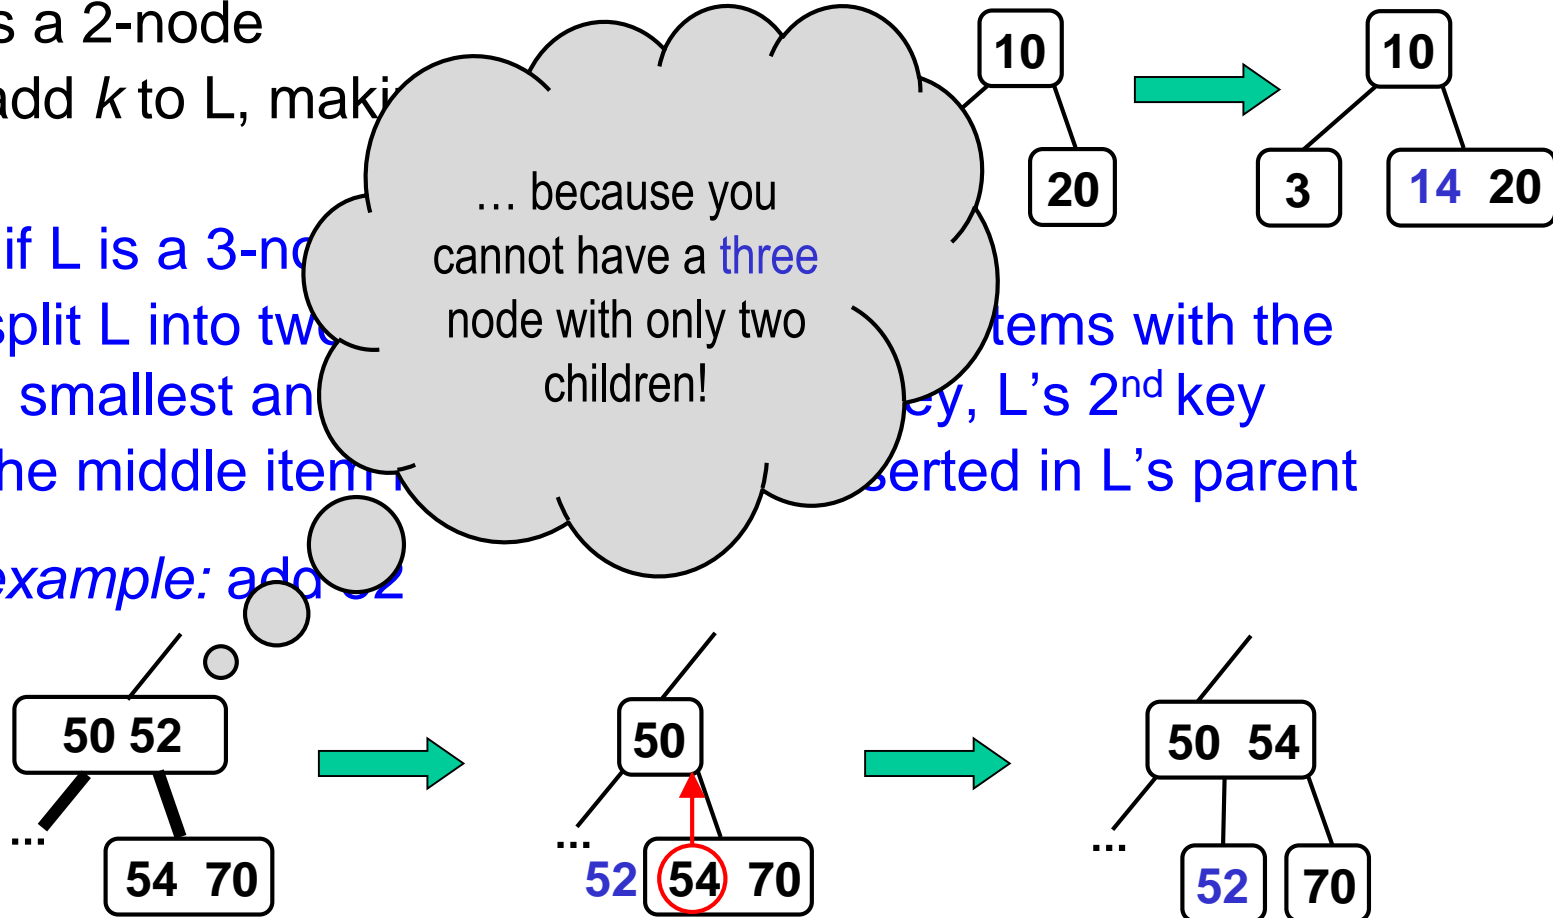
- Algorithm for inserting an item with a key  $k$ :
  - search for  $k$ , but don't stop until you hit a leaf node
  - let  $L$  be the leaf node at the end of the search
  - if  $L$  is a 2-node
    - add  $k$  to  $L$ , making it a 3-node

else if  $L$  is a 3-node

- split  $L$  into two 2-nodes
- the smallest and largest items with the same key,  $L$ 's 2<sup>nd</sup> key, are promoted to  $L$ 's parent
- the middle item is inserted in  $L$ 's parent

example: add 52

... because you cannot have a three node with only two children!




# Insertion in 2-3 Trees

- Algorithm for inserting an item with a key  $k$ :
  - search for  $k$ , but don't stop until you hit a leaf node
  - let  $L$  be the leaf node at the end of the search
  - if  $L$  is a 2-node
    - add  $k$  to  $L$ , making it a 3-node

else if  $L$  is a 3-node

- split  $L$  into two 2-nodes
- the smallest and largest keys remain with the original 2-node
- the middle item is "sent up" and inserted in  $L$ 's parent

example: add 52






# Insertion in 2-3 Trees

- Algorithm for inserting an item with a key  $k$ :
  - search for  $k$ , but don't stop until you hit a leaf node
  - let  $L$  be the leaf node at the end of the search
  - if  $L$  is a 2-node
    - add  $k$  to  $L$ , making it a 3-node

else if  $L$  is a 3-node


- split  $L$  into two 2-nodes
- the smallest and largest keys from  $L$  move up to become the new parent's keys, with the middle item is "sent up" and inserted in  $L$ 's parent

*example: add 52*




## Example 1: Insert 8

- Search for 8:




- Add 8 to the leaf node, *making it a 3-node*:




## Example 2: Insert 17


- Search for 17:



- Split the leaf node, and send up the middle of 14, 17, 20 and insert it the leaf node's parent:




## Example 3: Insert 92




## Example 3: Insert 92

- Search for 92:




- Split the leaf node, and send up the middle of 92, 93, 97 and insert it the leaf node's parent:




- In this case, the leaf node's parent is also a 3-node, so we need to split it as well...

## Example 3 (cont.)

- We split the root's right child, but the root is also full!




- Then we split the root, which increases the tree's height by 1, but the tree is still balanced.




- This is only case in which the tree's height increases.

# Efficiency of 2-3 Trees



- A 2-3 tree containing  $n$  items has a height  $\leq \log_2 n$ .
- Thus, search and insertion are both  $O(\log n)$ .
  - a search visits at most  $\log_2 n$  nodes
  - an **insertion** begins with a search; in the worst case, it goes all the way back up to the root performing splits, so it visits at most  $2\log_2 n$  nodes

# Efficiency of 2-3 Trees



- A 2-3 tree containing  $n$  items has a height  $\leq \log_2 n$ .
- Thus, search and insertion are both  $O(\log n)$ .
  - a search visits at most  $\log_2 n$  nodes
  - an insertion begins with a search; in the worst case, it goes all the way back up to the root performing splits, so it visits at most  $2\log_2 n$  nodes
- Deletion is tricky – you may need to coalesce nodes! However, it also has a time complexity of  $O(\log n)$ .
- Thus, we can use 2-3 trees for a  $O(\log n)$ -time data dictionary!