

SLS Lecture 8 : Writing some simple assembly programs

Contents

- [8.1. Writing a `popcnt` assembly program](#)
- [8.2. Writing an `add` assembly program](#)
- [8.3. Ending / Exiting our Program/Process](#)
- [8.4. Exercises and extra materials](#)

Spend some time writing some very simple assembly programs and learn to use the debugger so that we have enough skills to explore how things work. We will be repeat various things in more detail in future lectures.

- Write `popcnt` in assemble code
 - use gdb to play with the popcnt program
- Write a simple `add` in assembly code
 - use gdb to play with the add program
 - using the cpu as a glorified calculator
 - first pass at CPU support for “numbers”
- What happens if we let our programs continue
 - how do we successfully “halt/end” our execution
 - `int3` trap
 - tells OS to return control to debugger
 - more generally how can we make a Kernel/System Call
 - revisit `add` programs adding exits
 - `int3`
 - `exit` syscall
- Implicitly use our shell, editor, Make and Git knowledge to do the above

8.1. Writing a `popcnt` assembly program

- Write a one instruction assembly program
 1. first using `.byte`
 2. using intel assembly instruction
- Use gdb to explore how this instruction works
 - learn to use gdb to set register values
 - and how to execute and re-execute an instruction

8.1.1. Setup

1. make directory
2. open emacs and write `popcnt.s`

Skipping `git` for time.

1. make a directory for our work :
 - `mkdir simpleasm`
 - `cd simpleasm`
2. emacs popcnt.s

CODE: asm - The 'popcnt' assembly program

```
.intel_syntax noprefix

.section .text

.global _start
_start:
    popcnt rax, rbx    # same as .byte 0xF3, 0x48, 0x0F, 0xB8, 0xC3
```

Here is a fully commented version of the same code.

CODE: asm - The commented 'popcnt' assembly program

```
/* General anatomy of a assembly program line
[label]: <directive or opcode> [operands] # comment
*/

.intel_syntax noprefix # assembler syntax to use <directive>
                        # set assembly language format to intel

.section .text          # linker section <directive>
                        # let the linker know that what follows are cpu instructions to
                        # to be executed -- uposed to values that represent data.
                        # For historical reasons cpu instructions are called "text"

.global _start          # linker symbol type <directive>
                        # makes the symbol _start in this case visible to the linker
                        # The linker looks for an _start symbol so that it knows address
                        # of the first instruction of our program

_start:                # introduce a symbolic (human readable) label for "this" address
                        # associates the address of this point in our program with the
                        # name following the ':' -- in our case _start
                        # In our program or in the debugger we can use this name to
                        # to refer to this location -- address. And thus the values
                        # that end up here.

# .byte 0xF3, 0x48, 0x0F, 0xB8, 0xC3 # popcnt rax,rbx
popcnt rax, rbx        # ok the single intel opcode (instruction) that makes up
                        # our program

/*
    Details about the assembler directives and general syntax that we will be using
    https://sourceware.org/binutils/docs/as/
    Intel instruction set reference -- documents the cpu memonics/instructions of
    the computer's processor that we are writing for
    https://software.intel.com/content/www/us/en/develop/download/intel-64-and-ia-32-architectures-sdm-combined-volumes-2a-2b-2c-and-2d-instruction-set
*/
```

We can use the `.byte` directive to set the values in memory to anything we like eg.

```
.byte 0xF3, 0x48, 0x0F, 0xB8, 0xC3
```

But of course the real value is that we could have also simply written

```
popcnt rax, rbx
```

```
#To assemble and link the code we will use the following command:
as -g popcnt.s -o popcnt.o → object
ld -g popcnt.o -o popcnt → linker

# We can automate this using a makefile so that all we would need to do is:
make popcnt
```

```
# To get the debugger going:
gdb -x setup.gdb popcnt

# setup.gdb set intel assembly syntax and configures tui

# Now we want use gdb command to poke around popcnt

# Set a breakpoint at the start symbol so exection will stop their:
break _start
# Start the program running:
run
# play around with popcnt

set $rbx = 0b11 + set $PC = _start → rip
set $rbx = 0xFFFF
```

8.2. Writing an add assembly program

- re-enforce the steps to creating and debugging an assembly program
 - begin to explore CPU support for working with “numbers”
 - we will get into how numbers “work” later
 - learn enough so that you can poke around yourself
 - get an idea of cool things that INTEL instructions can do
 - try adding some variables in memory to our program
- Lets work with the `add` instruction in a similar way that we did with `popcnt`
- explore the results of adding with binary, hex, unsigned and signed values
- explore overflow *oflags → show flag*
- then make the program a little more complex:

```
movabs rbx, 0xdeadbeefdeadbeef
mov rax, 1
add rax, rbx
```

- lets use some more cool features of the intel instruction set

```
rdrand rbx
mov rax, 1
add rax, rbx
popcnt rbx, rax
```

- lets get a brief glimpse at how to use memory locations for the value

```
.intel_syntax noprefix
.data → variable memory set (static memory)
x: data ← .quad 142
y: time ← .quad 4200
sum: .quad

.text
.global _start
_start:
mov rax, QWORD PTR x
add rax, QWORD PTR y
mov QWORD PTR sum, rax
```

```
.intel_syntax noprefix

.text
.global _start
_start:
    add rax, rbx
```

```
# add targets to Makefile
make add

as -g add.s -o add.o
ld -g add.o -o add
```

```
gdb -x setup.gdb add
break _start
run
```

8.2.1. Exercises

- try repeating what we did with **add** with imul, and, or, xor: for each
 - create a new file
 - add targets to Makefile for it
 - use gdb to explore what the instruction does

8.3. Ending / Exiting our Program/Process

- What happens if we run our programs outside of the debugger?
 - why does this happen?

8.3.1. How can we avoid this *2 ways*

1. TRAP: Use an instruction that tells the OS to

- stop the process and give control back to the debugger
- if no debugger is running just kill process and signal shell
 - Instruction: `int3`:
 - Opcode: `0xCC`
 - Description: Interrupt 3 – trap to debugger

2. Call OS Kernel Exit Process call

- This is an example of calling an OS Kernel call to have the kernel do something for your process
- We will look at this more but for the moment here is what is necessary to call `exit`
 - pass return value to Kernel
 - exit/terminate process

8.3.2. Interrupt 3 `int3` – trap to debugger

64-ia-32-architecture...
Page 559 of 2,198

INSTRUCTION SET REFERENCE, A-L

INT *n*/INTO/INT 3—Call to Interrupt Procedure

Opcode	Instruction	Op/En	64-Bit Mode	Compat/Leg Mode	Description
CC	INT 3	NP	Valid	Valid	Interrupt 3—trap to debugger.
CD <i>ib</i>	INT <i>imm8</i>	I	Valid	Valid	Interrupt vector specified by immediate byte.
CE	INTO	NP	Invalid	Valid	Interrupt 4—if overflow flag is 1.

Instruction Operand Encoding

Op/En	Operand 1	Operand 2	Operand 3	Operand 4
NP	NA	NA	NA	NA
I	<i>imm8</i>	NA	NA	NA

Description

The INT *n* instruction generates a call to the interrupt or exception handler specified with the destination operand (see the section titled “Interrupts and Exceptions” in Chapter 6 of the *Intel® 64 and IA-32 Architectures Software Developer’s Manual, Volume 1*). The destination operand specifies a vector from 0 to 255, encoded as an 8-bit unsigned intermediate value. Each vector provides an index to a gate descriptor in the IDT. The first 32 vectors are reserved by Intel for system use. Some of these vectors are used for internally generated exceptions.

The INT *n* instruction is the general mnemonic for executing a software-generated call to an interrupt handler. The INTO instruction is a special mnemonic for calling overflow exception (#OF), exception 4. The overflow interrupt checks the OF flag in the EFLAGS register and calls the overflow interrupt handler if the OF flag is set to 1. (The INTO instruction cannot be used in 64-bit mode.)

The INT 3 instruction generates a special one byte opcode (CC) that is intended for calling the debug exception handler. (This one byte form is valuable because it can be used to replace the first byte of any instruction with a breakpoint, including other one byte instructions, without over-writing other code). To further support its function as a debug breakpoint, the interrupt generated with the CC opcode also differs from the regular software interrupts as follows:

- Interrupt redirection does not happen when in VME mode; the interrupt is handled by a protected-mode handler.
- The virtual-8086 mode IOPL checks do not occur. The interrupt is taken without faulting at any IOPL level.

Note that the “normal” 2-byte opcode for INT 3 (CD03) does not have these special features. Intel and Microsoft assemblers will not generate the CD03 opcode from any mnemonic, but this opcode can be created by direct numeric code definition or by self-modifying code.

The action of the INT *n* instruction (including the INTO and INT 3 instructions) is similar to that of a far call made with the CALL instruction. The primary difference is that with the INT *n* instruction, the EFLAGS register is pushed onto the stack before the return address. (The return address is a far address consisting of the current values of the CS and EIP registers.) Returns from interrupt procedures are handled with the IRET instruction, which pops the EFLAGS information and return address from the stack.

The vector specifies an interrupt descriptor in the interrupt descriptor table (IDT); that is, it provides index into the IDT. The selected interrupt descriptor in turn contains a pointer to an interrupt or exception handler procedure. In protected mode, the IDT contains an array of 8-byte descriptors, each of which is an interrupt gate, trap gate, or task gate. In real-address mode, the IDT is an array of 4-byte far pointers (2-byte code segment selector and a 2-byte instruction pointer), each of which point directly to a procedure in the selected segment. (Note that in real-address mode, the IDT is called the **interrupt vector table**, and its pointers are called interrupt vectors.)

The following decision table indicates which action in the lower portion of the table is taken given the conditions in the upper portion of the table. Each Y in the lower section of the decision table represents a procedure defined in the “Operation” section for this instruction (except #GP).

INT *n*/INTO/INT 3—Call to Interrupt Procedure

Vol. 2A 3-457

```
.intel_syntax noprefix

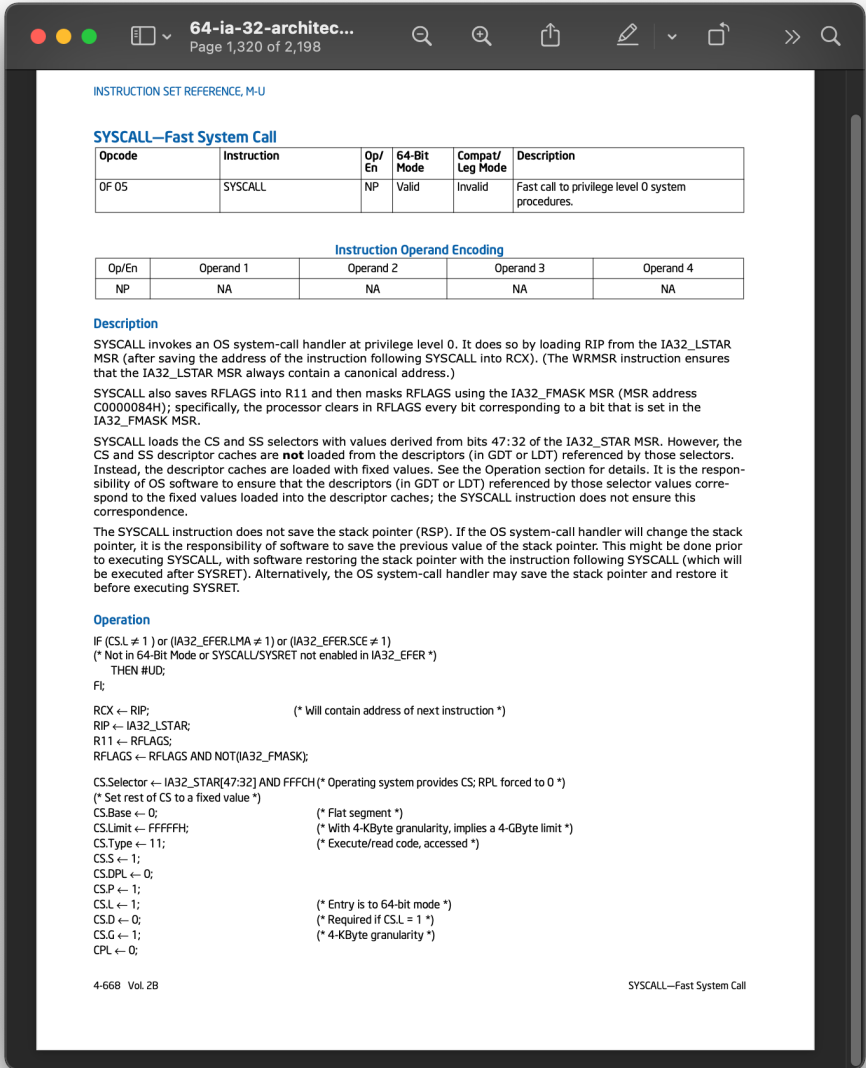
.text
.global _start
_start:
    int3
```

8.3.3. Exit – An OS service to terminate a process

To exit your process and return an exit value

- requires a call to the OS!

On Intel the instruction is `syscall`



8.3.4. The OS System Calls

Each OS Kernel provides a set of calls that an process can invoke using the `syscall` instruction on an Intel based computer

The Linux Kernel supports a very large number of system calls each is identified by a unique number that must be placed in `RAX` prior to executing the `syscall` instruction. Additional arguments are passed in by setting other registers.

With each version of the Kernel the table of calls changes. Here is one site that provides a list

Filippo.io

Searchable Linux Syscall Table for x86 and x86_64

There are some tables like this around, but they are usually cool auto-generated hacks and that has the downfall of the different implementations is the correct one, etc.

So, here is a lovingly hand-crafted Linux Syscall table for the x86[-64] architecture, with arguments, calling code included. Also, fuzzy search!

64-bit 32-bit
(Coming soc

Instruction: `syscall`
Return value found in: `%rax`
 Syscalls are implemented in functions named as in the *Entry point* column, or with the `DEFINE_SYSCALLx(%nam`
 Relevant man pages: [syscall\(2\)](#), [syscalls\(2\)](#).
 Double click on a row to reveal the arguments list. Search using the fuzzy filter box.

Filter:

- From the above we can see that the `exit` system call number is `60`
- reading some man pages `man syscall` and `man syscalls` we find that
 - we must place `60` in `rax`
 - and that the value we want to return in `rdi`

```

.intel_syntax noprefix

.text
.global _start
_start:
mov    rax, 60    # Linux exit system call number is 60
mov    rdi, 0     # rdi is return value 0 success

syscall

```

We will revisit OS system calls in more detail later

- this is good enough for the moment

8.3.4.1. Avoiding Hard coding system call numbers

Operating system code usually provides files that you can include in your code so that you don't have to hardcode magic numbers like 60 for exit. In Linux you can add the following file `#include <asm/unistd_64.h>` to get all the system call numbers. You can then use `__NR_exit` to mean the number for the exit system call.

eg. `exitfancy.S`

```

#include <asm/unistd_64.h>
.intel_syntax noprefix
.text
.global _start
_start:
mov    rax, __NR_exit # exit system call number
mov    rdi, 0         # UNIX success value is 0
syscall              # call OS. This will not return

```

replace this to 60

But the assemble does not support have support for including files. We must first use another tool called a preprocessor eg.

```

cc -E exitfancy.S > exitfancy.s
as -g exitfancy.s -o exitfancy.o
ld -g exitfancy.o -o exitfancy

```

In general we will just skip this and we will just use hardcoded numbers.

8.4. Exercises and extra materials

- rewrite all the examples to use `int3` at the end
- rewrite all the examples to call OS exit call
- combine some of the examples
- see what happens when you add 1 to `0xffffffffffff` using the add instructions.
 - any idea what is going on?

8.4.1. Makefile for all the lecture examples

```

popcnt: popcnt.o
ld -g popcnt.o -o popcnt

popcnt.o: popcnt.s
as -g popcnt.s -o popcnt.o

add: add.o
ld -g add.o -o add

add.o: add.s
as -g add.s -o add.o

exit: exit.o
ld -g exit.o -o exit

exit.o: exit.s
as -g exit.s -o exit.o

int3: int3.o
ld -g int3.o -o int3

int3.o: int3.s
as -g int3.s -o int3.o

exitfancy: exitfancy.o
ld -g exitfancy.o -o exitfancy

exitfancy.o: exitfancy.s
as -g exitfancy.s -o exitfancy.o

exitfancy.s: exitfancy.S
cc -E exitfancy.S > exitfancy.s

clean:
rm -f $(wildcard *.o popcnt add int3 exit exitfancy exitfancy.s)

```

8.4.2. Here is a fully documented fancy version of exit

- We use the the preprocessor to include the OS system call numbers
- and we use the `.equ` directive of the assembler to make our code more readable

A commented version that avoids "magic" numbers.

```

# Pull in a file that contains all the OS system call numbers
# for this to work we must preprocess this file via gcc -E before
# we can assemble it with as so that the contents of the header file
# will be included

#include <asm/unistd_64.h>

# See discussion bellow to understand what is in this file

.intel_syntax noprefix # set assembly language format to intel

# Define some constants so that we don't have magic numbers
# in our code. We use the .equ (equal) assembly directive
# "This directive sets the value of symbol to expression."
# (https://sourceware.org/binutils/docs/as/Equ.html#Equ)
# FORMAT:
# .equ <SYMBOL>, <EXPRESSION>

.equ UNIX_SUCCESS_VALUE, 0
.equ LINUX_SYSCALL_EXIT, __NR_exit

.text # Place the following in the area that
# instruction should be encoded and stored
# for historical reasons it is called text

.global _start

_start: # The linker knows to mark the _start address
# as location where execution should begin.
# The OS will be sure to setup the CPU so that the
# program counter is initialized with this address

# To voluntarily hand control of the cpu back to the Operating system so it can
# do something for us we use a special instruction
# -- on x86_64 this instruction is syscall (man syscall - table 1)
#
# Operating systems provide many functions that a program can call to
# get things done. eg. open a stream, read or write bytes to or from a stream
# or simply terminate and exit our program.
#
# Each call is identified by a number we call a system call number.
# On LINUX it expects us to load rax register with the system call number
# prior to executing the syscall. Once control is handed over to the OS
# it will then look the number up and invoke the appropriate function inside itself.
# Other registers can be used to pass parameters to the particular OS system call function
# that we want to invoke. The following is the conventions on X86_64 Linux
#
# Arch/ABI  arg1 arg2 arg3 arg4 arg5 arg6 arg7
# -----
# x86-64    rdi rsi rdx r10 r8  r9  -
# (see man syscall -- table 2)
# The list of all Linux system call numbers can be found in a file so that programmers
# can know what number to use to invoke a particular call
# On our installation of linux you can find it here
# /usr/include/x86_64-linux-gnu/asm/unistd_64.h
# For an explanation of all the systems calls see man syscalls (note the s)
# -- When we switch to understanding C we it will be easier to understand
# the above manual pages
#
# On LINUX
# One of these calls is the EXIT system call that tells OS our program is done
# (see man 2 exit)

mov    rax, LINUX_SYSCALL_EXIT # load rax with the Linux system call number for
# exit

mov    rdi, UNIX_SUCCESS_VALUE # load rdi with the exit status value we want
# to pass back to the program that launched us
# eg. the shell in which we ran our program from

syscall # use the syscall instruction to hand control over
# to the OS. Note since we will be terminated
# this instruction will be our last and nothing else
# need follow it

```