

UC-SLS Lecture 18 : Using C to write and organize opcode bytes - Functions

Contents

- 18.1. Opcodes and C
- 18.2. An example: Trace this code and visualize the stack

- create a directory `mkdir cfuncs; cd cfuncs`
- copy examples
- add a `Makefile` to automate assembling and linking
 - we are going run the commands by hand this time to highlight the details
- add our `setup.gdb` to make working in gdb easier
- normally you would want to track everything in git

```
$ ls /home/jovyan/cfuncs
Makefile    main0.c    myadd.c    myfunc0.c    myfunc2.c    setup.gdb
callmyadd.c  main0e.c    myadd.h    myfunc1.c    myfunc3.c
$
```

18.1. Opcodes and C

When we write assembly code we are free to layout our opcodes and use registers in any way we like.

- We can place labels anywhere in our opcodes
- We can specify a jump to any arbitrary location
- While we can use processor support for passing return address via instructions, like `call` and `ret`, we are not required too
- We are not forced to use the registers in any particular way

18.1.1. C Standardizes how to organize and write opcodes

- Its all about standardizing how things are done
 - https://www.uclibc.org/docs/psABI-x86_64.pdf
 - this way code written by different people or tools can inter-operate
 - there are rules they can rely on
- "C" forces us to decompose and organize opcodes into "functions"
 - global label - single entry point
 - block of opcodes ending in a "return"
 - standardizes use of registers
 - standardizes use of stack
 - `call` frames - automatic storage of locals
 - separation into declaration (many) and definition (one)
 - compiler can get smart and optimize functions and variables away
 - in-line
 - dead-code elimination
 - register only variables

18.1.2. C Standardizes how to organize and write opcodes

Overall summary

- "C" forces us to decompose and organize opcodes into "functions"
- "C" functions:
 - Have a unique global label that identifies a single starting address for the function
 - the label is formed from the function's name which must conform to certain rules
 - Form a contiguous block of memory that does not overlap with another function or data
 - Therefore they have a clear size in bytes that spans their first opcode to the last
 - Are written so that there is a standard way for passing arguments to them
 - a fixed way for passing arguments eg. basics on x86_64
 - `arg1 → rdi`
 - `arg2 → rsi`
 - `arg3 → rdx`
 - `arg4 → r8`
 - `arg5 → r9`
 - rest are pushed on stack in reverse order (`arg6` is last to be pushed)
 - a fixed way for passing a return value eg. basics on x86_64
 - return value → `rax`
 - details for cpu and OS are in specification documents eg:
 - https://www.uclibc.org/docs/psABI-x86_64.pdf
 - Execution from a function must end with a return to the next instruction after the `call`
 - eg on x86_64 `call` and `ret` are used
 - Support local variables that are automatically managed
 - eg. on x86_64 the processor stack is used
 - each `call` to a function creates a new stack frame
 - each frame represents a `call` to a function
 - the frame contains a version of the local variables for that call
 - thus each `call` has its own locals
 - when a function returns to its `call` the stack frame for the `call` is popped

- thus support recursion
- Separates declaration from definition
 - declaration: only specify its name but does not define any opcodes
 - the function declaration is needed to generate the assembly code to call the function
 - given the rules above for how arguments are passed and a unique entry point
 - compiler can generate the assembly code for a call with the declaration
 - does not need the body
 - declarations are placed in "header" files.
 - definitions are placed in "c" files.
 - definition: repeats the declaration but include a body that defines the functions opcodes
 - there can only be one of these
- A compiler is allowed to "inline" a function if certain optimizations are enabled and criteria are met
 - there are times in which the overhead of calling a function is not worth it
 - it is better to just create a version of the opcodes in place where the call is being made
 - inlined

18.1.3. Let's start at the beginning

- we will use the compiler and our ability to read assembly code to learn how "C" works

18.1.3.1. Function Name → global label for its Entry point (start of its opcode)

C:

```
void myfunc(void) {}
```

```
gcc -fno-inline -fno-stack-protector -fno-pic -static -Werror -fcf-protection=none -fno-asynchronous-unwind-tables -Os -S -masm=intel myfunc0.c -o myfunc0.s
```

Assembly:

```
.file "myfunc0.c"
.intel_syntax noprefix
.text
.globl myfunc
.type myfunc, @function
myfunc:
    .ret
    .size myfunc, .-myfunc
    .ident "GCC: (Ubuntu 9.4.0-lubuntu1~20.04.1) 9.4.0"
    .section .note.GNU-stack,"@progbits"
```

- function name "myfunc" introduces a global label **myfunc** in the text section
- return type prefixes function **void** is the "no" type
 - a **void** return type means that the function does not return anything
- parenthesis after function name (and) demarks parameter list
 - a **void** in parameter list means function takes no parameters
- { and } demarks body
 - a set of statements that will be converted into opcodes
 - implicitly every function has at least one statement
 - **return**
 - if not written the compiler assumes one
 - generates instructions to return to the caller
 - X86: **ret**

18.1.3.2. Calling a function

C:

```
__attribute__((noinline))
int funcA(void) {
    return 7;
}

int funcB(void) {
    return 3 + funcA();
}
```

```
gcc -fno-inline -fno-stack-protector -fno-pic -static -Werror -fcf-protection=none -fno-asynchronous-unwind-tables -Os -S -masm=intel myfunc1.c -o myfunc1.s
```

Assembly

```

.file  "myfunc1.c"
.intel_syntax noprefix
.text
.globl funcA
.type  funcA, @function
funcA:
    mov    eax, 7
    ret
.size  funcA, .-funcA
.globl funcB
.type  funcB, @function
funcB:
    call   funcA
    add    eax, 3
    ret
.size  funcB, .-funcB
.ident "GCC: (Ubuntu 9.4.0-1ubuntu1~20.04.1) 9.4.0"
.section .note.GNU-stack,"",@progbits

```

- avoid compiler for optimizing away `noinline` and use of return value
- exactly what we expected in assembly right?

18.1.4. Interlude : `_start` vs `main()`

- As we have seen the linker marks where execution should begin in our binary via the `_start` symbol
- However when we are writing 'C' we normally do not write raw assembly
 - all our code is in functions
- In the last lectures we wrote our own `_start` in an assembly file that called our C generated assembly
 - and linked it by hand avoid all the defaults
- However normally we don't do this.
 - the C compiler come with some startup code along with the standard C library of functions
 - this code is usually in an object files of the name `crt*.o`
 - The "c" runtime a bunch of code that runs before the code you write
 - runs setup code for you initializing c library and other aspects
 - when done calls `main` function passing in some standard parameters
 - `argc`, `argv` and on Unix `envp`
 - use `-v` to see it get added
 - So lets write a main and use gcc to link it in the "normal" way

C:

```

int main() {
    myfunc();
    return 0;
}

```

```

gcc -fno-inline -fno-stack-protector -fno-pic -static -Werror -fcf-
protection=none -fno-asynchronous-unwind-tables -Os -S -masm=intel main0e.c -o
main0e.s
main0e.c: In function 'main':
main0e.c:2:3: error: implicit declaration of function 'myfunc' [-Werror=implicit-
function-declaration]
  2 |     myfunc();
     |     ^~~~~~
ccl: all warnings being treated as errors
make[2]: *** [Makefile:34: main0e.s] Error 1

```

What's the fix?

18.1.4.1. Function Declarations vs Definitions

- When C encounters code that calls another function
 - it cannot know how to generate the assembly for the call
 - unless it knows
 - name
 - arguments: type and order
 - and return value type
 - a function declaration is exactly this - just the signature of the function
 - does not generate any code itself just allows calls to be correctly created
 - a single definition either in the same file or another that will end up in a .o must exist
 - linker will stitch them together

Add a declaration of `myfunc` that matches its defintion to the `main.c` file

Normally we would put this into a header file eg. `myfunc0.h`

- this way any file in which we want to call `myfunc` in we would simply include the header

C:

```

void myfunc(void);

int main() {
    myfunc();
    return 0;
}

```

```

gcc -fno-inline -fno-stack-protector -fno-pic -static -Werror -fcf-
protection=none -fno-asynchronous-unwind-tables -Os -S -masm=intel main0.c -o
main0.s

```

Assembly

```

.file "main0.c"
.intel_syntax noprefix
.text
.section .text.startup,"ax",@progbits
.globl main
.type main, @function
main:
    push rax
    call myfunc
    xor eax, eax
    pop rdx
    ret
.size main, .-main
.ident "GCC: (Ubuntu 9.4.0-1ubuntu1~20.04.1) 9.4.0"
.section .note.GNU-stack,"",@progbits

```

```
gcc -fno-pic -static -Os -g main0.s myfunc0.s -o main0
```

Lets add the verbose flag so that we can see what is really going on

```

gcc -v -fno-pic -static -Os -g main0.s myfunc0.s -o main0
Using built-in specs.
COLLECT_GCC=gcc
COLLECT_LTO_WRAPPER=/usr/lib/gcc/x86_64-linux-gnu/9/lto-wrapper
OFFLOAD_TARGET_NAMES=nvptx-none:hsa
OFFLOAD_TARGET_DEFAULT=1
Target: x86_64-linux-gnu
Configured with: ../src/configure -v --with-pkgversion='Ubuntu 9.4.0-
lubuntul~20.04.1' --with-bugurl=file:///usr/share/doc/gcc-9/README.Bugs --enable-
languages=c,ada,c++,go,brig,d,fortran,objc,obj-c++,gm2 --prefix=/usr --with-gcc-
major-version-only --program-suffix=-9 --program-prefix=x86_64-linux-gnu- --
enable-shared --enable-linker-build-id --libexecdir=/usr/lib --without-included-
gettext --enable-threads=posix --libdir=/usr/lib --enable-nls --enable-clocale=gnu
--enable-libstdcxx-debug --enable-libstdcxx-time=yes --with-default-libstdcxx-
abi=new --enable-gnu-unique-object --disable-vtable-verify --enable-plugin --
enable-default-pie --with-system-zlib --with-target-system-zlib=auto --enable-
objc-gc=auto --enable-multiarch --disable-werror --with-arch-32=i686 --with-
abi=m64 --with-multilib-list=m32,m64,mx32 --enable-multilib --with-tune=generic --
enable-offload-targets=nvptx-none:/build/gcc-9-Av3Ued/gcc-9-9.4.0/debian/tmp-
nvptx/usr,hsa --without-cuda-driver --enable-checking=release --build=x86_64-
linux-gnu --host=x86_64-linux-gnu --target=x86_64-linux-gnu
Thread model: posix
gcc version 9.4.0 (Ubuntu 9.4.0-1ubuntu1~20.04.1)
COLLECT_GCC_OPTIONS='-v' '-fno-pic' '-static' '-Os' '-g' '-o' 'main0' '-
mtune=generic' '-march=x86-64'
  as --gdwarf2 -v --64 -o /tmp/cCHKeGxJ.o main0.s
GNU assembler version 2.34 (x86_64-linux-gnu) using BFD version (GNU Binutils for
Ubuntu) 2.34
COLLECT_GCC_OPTIONS='-v' '-fno-pic' '-static' '-Os' '-g' '-o' 'main0' '-
mtune=generic' '-march=x86-64'
  as --gdwarf2 -v --64 -o /tmp/ccMnEWal.o myfunc0.s
GNU assembler version 2.34 (x86_64-linux-gnu) using BFD version (GNU Binutils for
Ubuntu) 2.34
COMPILER_PATH=/usr/lib/gcc/x86_64-linux-gnu/9/:/usr/lib/gcc/x86_64-linux-
gnu/9/:/usr/lib/gcc/x86_64-linux-gnu/:/usr/lib/gcc/x86_64-linux-
gnu/9/:/usr/lib/gcc/x86_64-linux-gnu/
LIBRARY_PATH=/usr/lib/gcc/x86_64-linux-gnu/9/:/usr/lib/gcc/x86_64-linux-
gnu/9/../../../../lib/x86_64-linux-gnu/:/lib/..../lib:/usr/lib/x86_64-linux-
gnu/:/usr/lib/..../lib:/usr/lib/gcc/x86_64-linux-gnu/9/../../../../lib/:/lib:/usr/lib/
COLLECT_GCC_OPTIONS='-v' '-fno-pic' '-static' '-Os' '-g' '-o' 'main0' '-
mtune=generic' '-march=x86-64'
  /usr/lib/gcc/x86_64-linux-gnu/9/collect2 -plugin /usr/lib/gcc/x86_64-linux-
gnu/9/liblto_plugin.so -plugin-opt=/usr/lib/gcc/x86_64-linux-gnu/9/lto-wrapper -
plugin-opt=fresolution=/tmp/ccZKBpMK.res -plugin-opt=pass-through=lgcc -plugin-
opt=pass-through=lgcc_eh -plugin-opt=pass-through=lgcc_eh --build-id=elf_x86_64
--hash-style=gnu --as-needed -static -relro -o main0 /usr/lib/gcc/x86_64-linux-
gnu/9/../../../../x86_64-linux-gnu/crti.o /usr/lib/gcc/x86_64-linux-gnu/9/crtbeginT.o
-L/usr/lib/gcc/x86_64-linux-gnu/9 -L/usr/lib/gcc/x86_64-linux-gnu/9/../../../../lib
-L/usr/x86_64-linux-gnu/9/Lib/..Lib -L/usr/lib/x86_64-linux-gnu -
L/usr/Lib/..Lib -L/usr/lib/gcc/x86_64-linux-gnu/9/../../../../tmp/cCHKeGxJ.o
/tmp/ccMnEWal.o --start-group -lgcc -lgcc_eh -lc --end-group /usr/lib/gcc/x86_64-
linux-gnu/9/crtend.o /usr/lib/gcc/x86_64-linux-gnu/9/../../../../x86_64-linux-
gnu/crti.o
COLLECT_GCC_OPTIONS='-v' '-fno-pic' '-static' '-Os' '-g' '-o' 'main0' '-
mtune=generic' '-march=x86-64'
```

Much bigger than we might have expected

```
-rwxr-xr-x. 1 joyyan root 853K Sep 29 16:39 main0
```

Why so big?

1. All the extra stuff
2. We have suppressed dynamic link/loading

Normally we just take all this on faith ... but since we know how now let's look at it at least just this once at what all this stuff is

```
objdump -d main0 > main0.dis
```

Edit

18.1.5. Variables declared in body of function are function local

- If the compiler needs memory for a local variable then it adds it to stack frame for call

C:

```

__attribute__((noinline))
void myfunc2(long long *i)
{
    *i += 1;
}

long long myfunc(void) {
    long long i=(long long)&myfunc;
    myfunc2(&i);
    return i;
}

```

```

gcc -fno-inline -fno-stack-protector -fno-pic -static -Werror -fcf-
protection=none -fno-asynchronous-unwind-tables -Os -S -masm=intel myfunc2.c -o
myfunc2.s

```

Assembly

```

.file   "myfunc2.c"
.intel_syntax noprefix
.text
.globl myfunc2
.type  myfunc2, @function
myfunc2:
    inc    QWORD PTR [rdi]
    ret
.size  myfunc2, .-myfunc2
.globl myfunc
.type  myfunc, @function
myfunc:
    sub   rsp, 16
    mov   QWORD PTR [rsp+8], OFFSET FLAT:myfunc
    lea   rdi, [rsp+8]
    call  myfunc2
    mov   rax, QWORD PTR [rsp+8]
    add   rsp, 16
    ret
.size  myfunc, .-myfunc
.ident "GCC: (Ubuntu 9.4.0-lubuntu1~20.04.1) 9.4.0"
.section .note.GNU-stack,"",@progbits

```

- must play some games to get the compiler to create a local with such simple functions
- notice what this code is doing?
- sub rsp, 16 what is this?
- what is rsp + 8
- Local variables have not fixed location in memory!
- PS this code is dangerous and not something you would normally do
 - but C lets you cheat if you want too!

18.1.6. LEA

LEA — SORRY THIS IS GOING TO HURT

▶ LOAD EFFECTIVE ADDRESS

▶ What does this do?

- ▶ MOVQ R8, [RAX + 2*RBX + 31]
 - ▶ cpu calculates Effective Address **ea** = RAX + 2 * RBX + 31
 - ▶ then loads R8 with 8 bytes starting at address **ea**
 - ▶ eg. loaded with little endian order with contents of M[ea] M[ea+1] ... M[ea+7]
- ▶ LEA R8, [RAX + 2*RBX + 31]
 - ▶ cpu calculates Effective Address **ea** = RAX + 2 * RBX + 31
 - ▶ the load R8 with **ea** – only time memory equation is used to calculate a number that you explicitly can use – DOES NOT FETCH VALUES FROM MEMORY!

If you are still confused by address modes you should now know enough that reading 3.7.5 "Specifying an Offset" of 3-22 Vol. 1 Intel SDM. should be a little easier. Also see 3-528 Vol. 2A – LEA–Load Effective Address

18.1.7. Passing arguments and simple return value

c:

```

__attribute__((noinline))
int func2(int x, int y)
{
    return x + y;
}

int func1(int x)
{
    return func2(x,2);
}

```

```

gcc -fno-inline -fno-stack-protector -fno-pic -static -Werror -fcf-
protection=none -fno-asynchronous-unwind-tables -Os -S -masm=intel myfunc3.c -o
myfunc3.s

```

Assembly

```

.file "myfunc3.c"
.intel_syntax noprefix
.text
.globl func2
.type func2, @function
func2:
    lea    eax, [rdi+rsi]
    ret
.size func2, .-func2
.globl func1
.type func1, @function
func1:
    mov    esi, 2
    jmp    func2
.size func1, .-func1
.ident "GCC: (Ubuntu 9.4.0-1ubuntu1~20.04.1) 9.4.0"
.section .note.GNU-stack,"",@progbits

```

- What is `lea` being used for here?
- Why `eax`, `rdi` and `rsi` (`esi`) being used the way they are?

LEA

- The compiler is smart and figure out that it can use `lea` for arbitrary math that does not require flags update

Why these Register?

CALLING CONVENTIONS FOR THE ISA

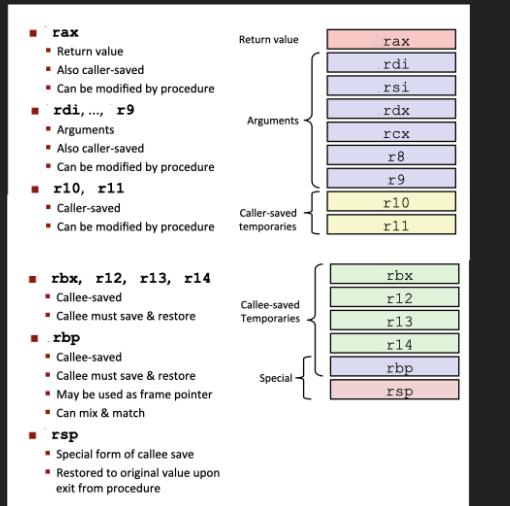
- ▶ Divide registers into two groups
 - ▶ Caller saves on stack prior to call if it is relying on value – Callee can blindly use as a temporary
 - ▶ Callee saves on frame before using so called can rely on the values in those registers to always be right

Register Saving Conventions

- When procedure you calls who:
 - you is the caller
 - who is the callee
- Can register be used for temporary storage?
- Conventions
 - “Caller Saved”
 - Caller saves temporary values in its frame before the call
 - “Callee Saved”
 - Callee saves temporary values in its frame before using
 - Callee restores them before returning to caller

CALLING CONVENTIONS FOR THE ISA

- ▶ Divide registers into two groups
 - ▶ Caller saves on stack prior to call if it is relying on value – Callee can blindly use as a temporary
 - ▶ Callee saves on frame before using so called can rely on the values in those registers to always be right
- ▶ Arguments:
 - ▶ Diane Sips Delicious Coffee 8 out of 9 times



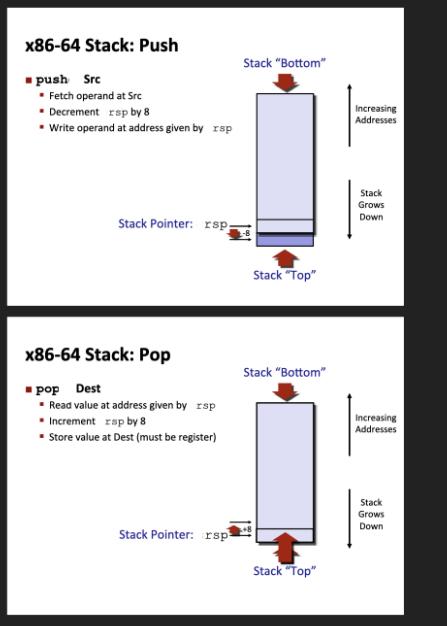
But remember the real truth or the rules for any given CPU and OS is in a standard somewhere

https://www.uclibc.org/docs/psABI-x86_64.pdf

- 3.2.3 Parameter Passing p17
- Figure 3.4: Register Usage p21

FUNCTIONS AND MEMORY FOR LOCAL VARIABLES — THE STACK AND FRAMES

- ▶ Using call and ret already make use of the stack so that we can transfer control and return back to and from a function
- ▶ C more formally defines how to use the stack for functions



FUNCTIONS AND MEMORY FOR LOCAL VARIABLES — THE STACK AND FRAMES

- ▶ Using call and ret already make use of the stack so that we can transfer control and return back to and from a function
- ▶ C more formally defines how to use the stack for functions – Stack oriented languages

Stack-Based Languages

- Languages that support recursion
 - e.g., C, Pascal, Java
 - Code must be “Reentrant”
 - Multiple simultaneous instantiations of single procedure
 - Need some place to store state of each instantiation
 - Arguments
 - Local variables
 - Return pointer
- Stack discipline
 - State for given procedure needed for limited time
 - From when called to when return
 - Callee returns before caller does
- Stack allocated in Frames
 - state for single procedure instantiation

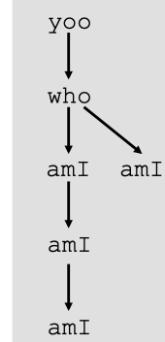
Call Chain Example

```
yoo (...)  
{  
    .  
    .  
    who ();  
    .  
    .  
}
```

```
who (...)  
{  
    . . .  
    amI ();  
    . . .  
    amI ();  
    . . .  
}
```

```
amI (...)  
{  
    .  
    .  
    amI ();  
    .  
    .  
}
```

Example
Call Chain

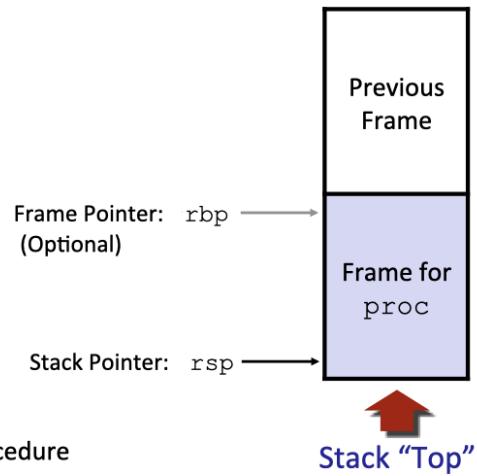


Procedure amI () is recursive

Stack Frames

■ Contents

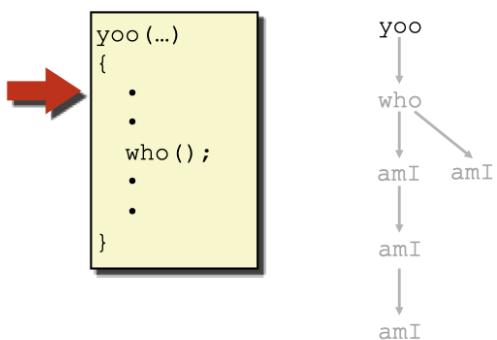
- Return information
- Local storage (if needed)
- Temporary space (if needed)



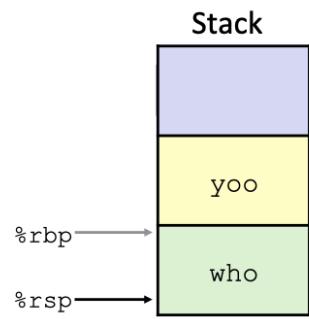
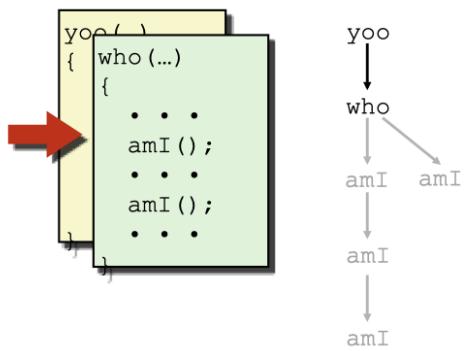
■ Management

- Space allocated when enter procedure
 - “Set-up” code
 - Includes push by **call** instruction
- Deallocated when return
 - “Finish” code
 - Includes pop by **ret** instruction

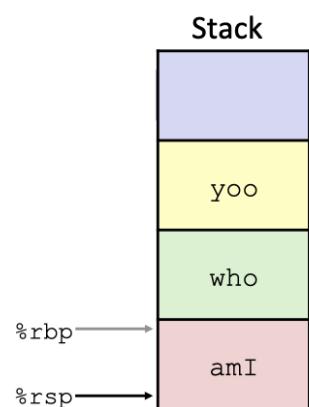
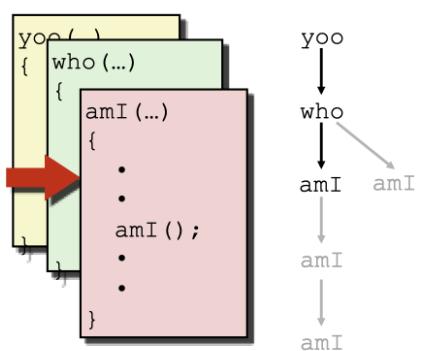
Example



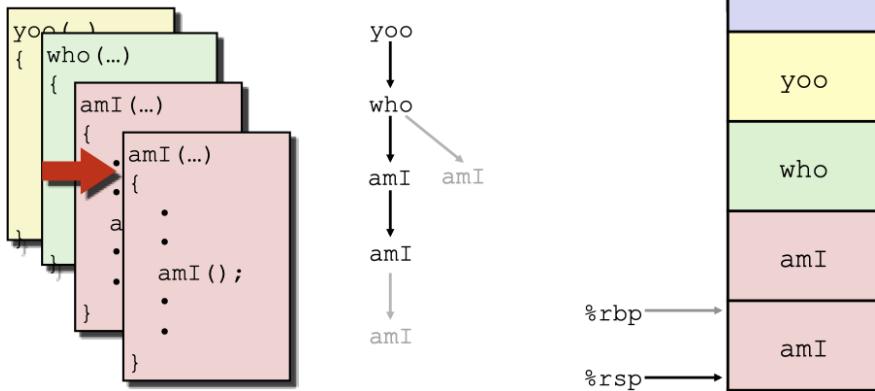
Example



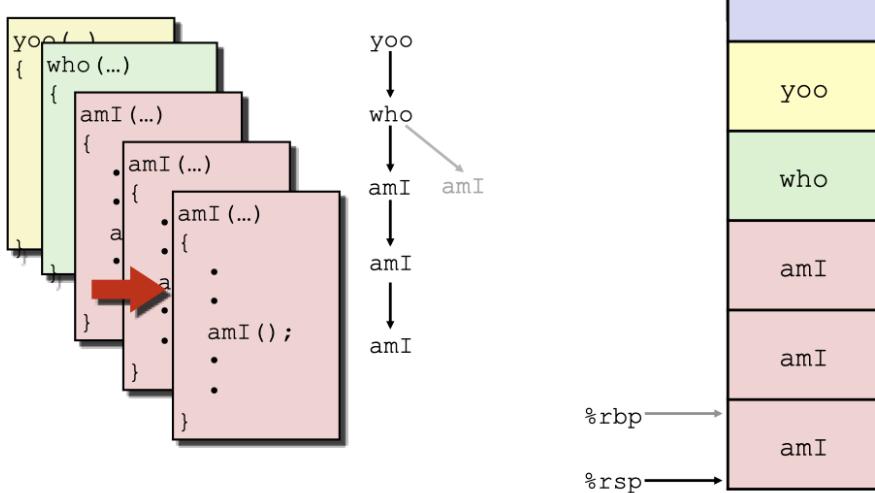
Example



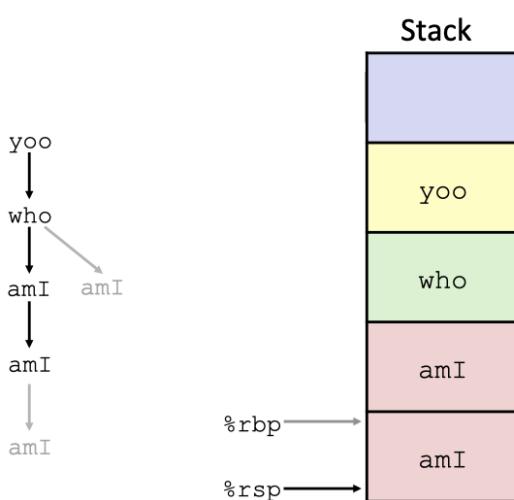
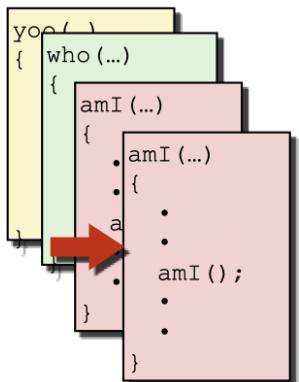
Example



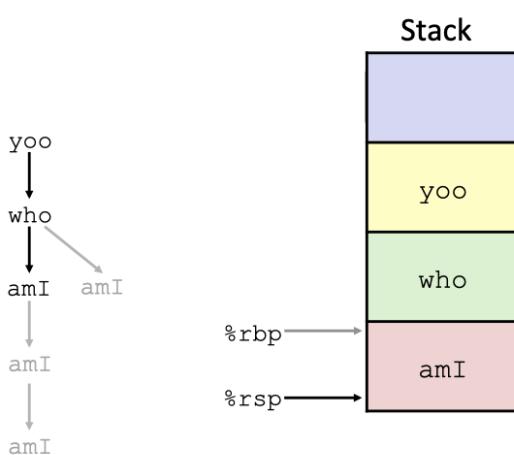
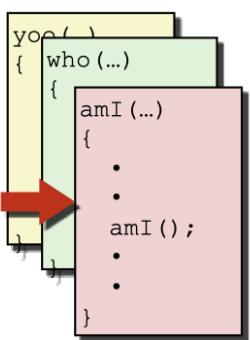
Example



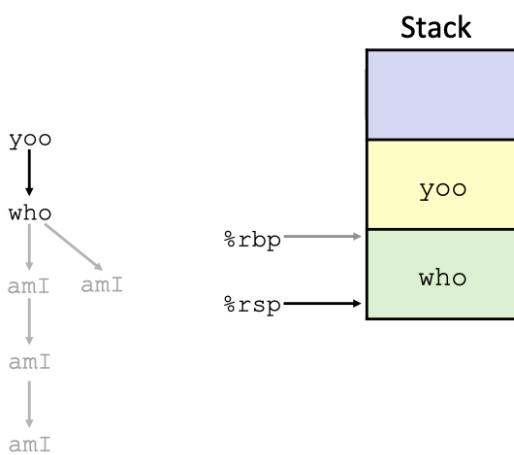
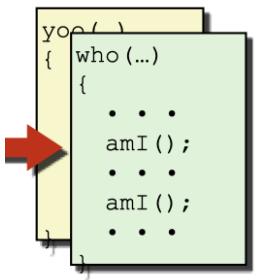
Example



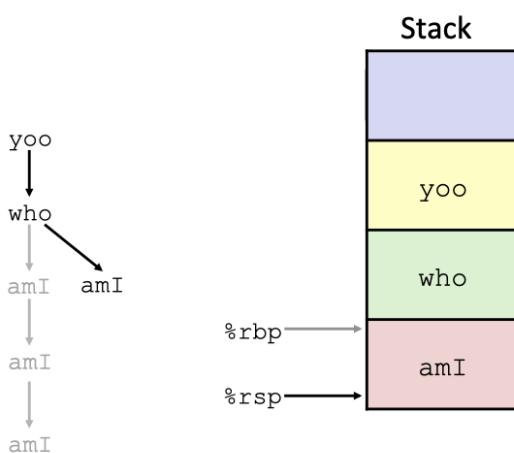
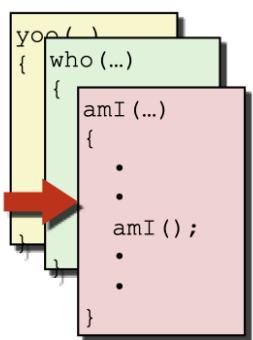
Example



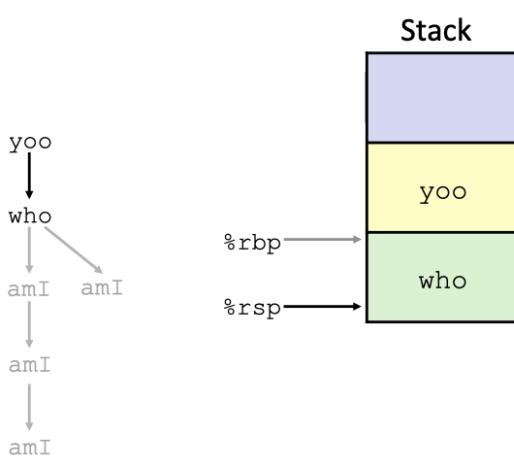
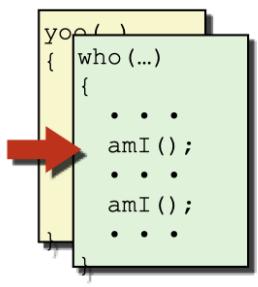
Example



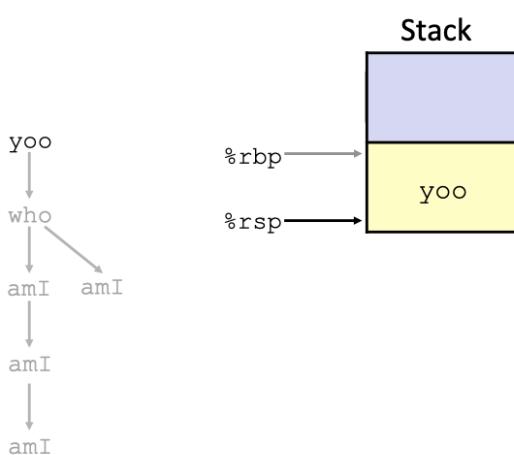
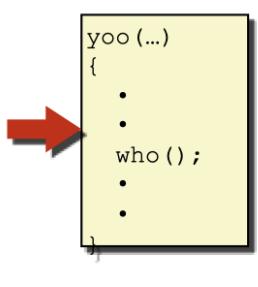
Example



Example



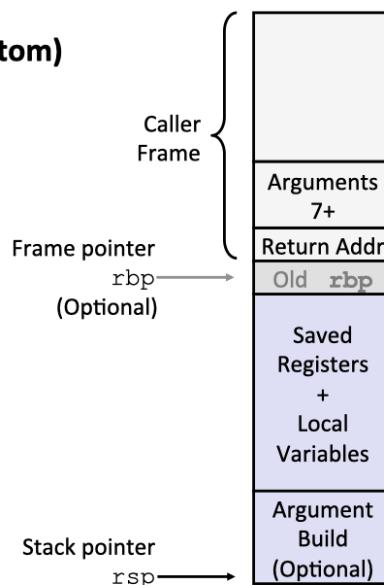
Example



x86-64/Linux Stack Frame

■ Current Stack Frame (“Top” to Bottom)

- “Argument build:”
Parameters for function about to call
- Local variables
If can’t keep in registers
- Saved register context
- Old frame pointer (optional)



■ Caller Stack Frame

- Return address
 - Pushed by `call` instruction
- Arguments for this call

18.2. An example: Trace this code and visualize the stack

C: myadd.c

```
long myadd(long *x_ptr, long val)
{
    long x = *x_ptr;
    long y = x + val;
    *x_ptr = y;
    return x;
}
```

```
gcc -fno-inline -fno-stack-protector -fno-pic -static -Werror -fcf-protection=none -fno-asynchronous-unwind-tables -Os -S -masm=intel myadd.c -o myadd.s
```

Assembly

```
.file "myadd.c"
.intel_syntax noprefix
.text
.globl myadd
.type myadd, @function
myadd:
    mov    rax, QWORD PTR [rdi]
    add    rsi, rax
    mov    QWORD PTR [rdi], rsi
    ret
.size myadd, .-myadd
.ident "GCC: (Ubuntu 9.4.0-1ubuntu1~20.04.1) 9.4.0"
.section .note.GNU-stack,"",@progbits
```

C: myadd.h

```
#ifndef __MY_ADD_H__
#define __MY_ADD_H__

long myadd(long *x_ptr, long val);

#endif // __MY_ADD_H__
```

C: callmyadd.c

```
#include "myadd.h"

long call_myadd(void) {
    long a = 15214;
    long b = myadd(&a, 5001);
    return a+b;
}

long
main() {
    return call_myadd();
}
```

```
gcc -fno-inline -fno-stack-protector -fno-pic -static -Werror -fcf-protection=none -fno-asynchronous-unwind-tables -Os -S -masm=intel callmyadd.c -o callmyadd.s
```

Assembly

```
.file "callmyadd.c"
.intel_syntax noprefix
.text
.globl call_myadd
.type call_myadd, @function
call_myadd:
    sub    rsp, 24
    mov    esi, 5001
    mov    QWORD PTR [rsp+8], 15214
    lea    rdi, [rsp+8]
    call   myadd
    add    rax, QWORD PTR [rsp+8]
    add    rsp, 24
    ret
.size  call_myadd, .-call_myadd
.section .text.startup,"ax",@progbits
.globl main
.type  main, @function
main:
    jmp   call_myadd
.size  main, .-main
.ident "GCC: (Ubuntu 9.4.0-1ubuntu1~20.04.1) 9.4.0"
.section .note.GNU-stack,"",@progbits
```