

C Functions

CS 210 - Fall 2023

Vasiliki Kalavri

vkalavri@bu.edu

Defining and Calling Functions

- Before we go over the formal rules for defining a function, let's look at our simple calculator program and how it defines functions.

return type

cor

```
// adds two integers
void add(int i, int j) {
    int res = i + j;
    printf("Result: %d\n", res);
}

// multiplies two numbers
int mul(int i, int j) {
    int k, res = 0;
    for (k=0; k<j; k++) {
        res += i;
    }
    printf("Result: %d\n", res);
    return res;
}
```

Return type

```
// adds two integers
void add(int i, int j) {
    int res = i + j;
    printf("Result: %d\n", res);
}

// multiplies two numbers
int mul(int i, int j) {
    int k, res = 0;
    for (k=0; k<j; k++) {
        res += i;
    }
    printf("Result: %d\n", res);
    return res;
}
```

Return type

```
// adds two integers
void add(int i, int j) {
    int res = i + j;
    printf("Result: %d\n", res);
}

// multiplies two numbers
int mul(int i, int j) {
    int k, res = 0;
    for (k=0; k<j; k++) {
        res += i;
    }
    printf("Result: %d\n", res);
    return res;
}
```

Parameters

```
int main() {  
    int i = 2;  
    power(2, i);  
    return 0;  
}  
  
// computes i^j using mul()  
int power(int i, int j) {  
    int res = 1;  
    while(j > 0) {  
        res = mul(res, i);  
        j--;  
    }  
    return res;  
}
```

```
int main() {  
    int i = 2;  
    power(2, i);  
    return 0;  
}
```

```
// computes i^j using mul()  
int power(int i, int j) {  
    int res = 1;  
    while(j > 0) {  
        res = mul(res, i);  
        j--;  
    }  
    return res;  
}
```

**Arguments are passed
by value**

- When a function is called, each argument is evaluated and its value assigned to the corresponding parameter.
- Since the parameter contains a copy of the argument's value, any changes made to the parameter during the execution of the function don't affect the argument.

```
int main() {  
    int i = 2;  
    power(2, i);  
    return 0;  
}  
  
// computes i^j using mul()  
int power(int i, int j) {  
    int res = 1;  
    while(j > 0) {  
        res = mul(res, i);  
        j--;  
    }  
    return res;  
}
```

Arguments are passed by value

Function Declarations

- C doesn't require that the definition of a function precede its calls.

```
#include <stdio.h>

int main() {

    int i, j;

    printf("Enter two numbers: ");
    scanf("%d %d", &i, &j);
    add(i, j);
    return 0;
}

// adds two integers
void add(int i, int j) {
    int res = i + j;
    printf("Result: %d\n", res);
}
```

compile time

Function Declarations

- C doesn't require that the definition of a function precede its calls.

```
#include <stdio.h>

int main() {

    int i, j;

    printf("Enter two numbers: ");
    scanf("%d %d", &i, &j);
    add(i, j);
    return 0;
}

// adds two integers
void add(int i, int j) {
    int res = i + j;
    printf("Result: %d\n", res);
}
```

```
$ gcc -fno-inline -fno-stack-protector -fno-pic -static -fcf-protection=none -fno-asynchronous-unwind-tables -Os -masm=intel -g add.c -o add
add.c: In function 'main':
add.c:7:3: warning: implicit declaration of function 'add' [-Wimplicit-function-declaration]
   7 |     add(i, j);
     |     ^~~
add.c: At top level:
add.c:12:6: warning: conflicting types for 'add'
   12 | void add(int i, int j) {
     |      ^~~
add.c:7:3: note: previous implicit declaration of 'add' was here
   7 |     add(i, j);
     |     ^~~
```

Function Declarations

- When the compiler encounters the first call of `add` in `main`, it has no information about the function.
- Instead of producing an error message, the compiler assumes that `add` returns an `int` value.
- The compiler has created an *implicit declaration* of the function.

Function Declarations

- A better solution: declare each function before calling it.
- A *function declaration* provides the compiler with a brief glimpse at a function whose full definition will appear later.
- General form of a function declaration:
return-type function-name (parameters) ;
- The declaration of a function must be consistent with the function's definition.

```
/**
 * A basic integer calculator.
 * It supports addition, subtraction,
 * multiplication, division, and power.
 *
 * **/

#include <stdio.h>
#include <stdlib.h>

void print_help();
void add(int, int);
int mul(int, int);
int division(int i, int j);
int power(int i, int j);
int rec_power(int i, int j);
```

```

int main() {

    int i, j;
    char op;

    // print instructions
    print_help();

    while (1) {
        printf("Enter operator or Q to exit: ");
        // read operation
        scanf(" %c ", &op);

        if (op != 'Q') {
            // read input numbers
            printf("Enter two integers: ");
            scanf("%d %d", &i, &j);

            if (i < 0 || j < 0) {
                printf("Negative numbers are ignored.\n");
                continue;
            }

            // apply operation
            switch (op) {
                case '+': add(i, j); break;
                case '-': printf("Result: %d\n", i-j); break;
                case '*': mul(i, j); break;
                case '/': {
                    int ret = division(i, j);
                    if (ret < 0) break;
                    printf("Result: %d\n", ret);
                    break;
                }
                case '^': printf("Result: %d\n", power(i, j)); break;
                default: printf("Invalid operator.\n"); break;
            }
        }
        else {
            printf("Sorry to see you go :/\n");
            return 0;
        }
    }
    return 0;
}

```

```

int main() {

    int i, j;
    char op;

    // print instructions
    print_help();

    while (1) {
        printf("Enter operator or Q to exit: ");
        // read operation
        scanf(" %c ", &op);

        if (op != 'Q') {
            // read input numbers
            printf("Enter two integers: ");
            scanf("%d %d", &i, &j);

            if (i < 0 || j < 0) {
                printf("Negative numbers are ignored.\n");
                continue;
            }

            // apply operation
            switch (op) {
                case '+': add(i, j); break;
                case '-': printf("Result: %d\n", i-j); break;
                case '*': mul(i, j); break;
                case '/': {
                    int ret = division(i, j);
                    if (ret < 0) break;
                    printf("Result: %d\n", ret);
                    break;
                }
                case '^': printf("Result: %d\n", power(i, j)); break;
                default: printf("Invalid operator.\n"); break;
            }
        }
        else {
            printf("Sorry to see you go :/\n");
            return 0;
        }
    }
    return 0;
}

```

```

/** Prints the calculator usage instructions */
void print_help() {
    printf("\n***Welcome to the basic calculator program!***\n\n");
    printf("Select one of the following operations or type 'Q' to exit\n");
    printf("+\t: Addition\n");
    printf("-\t: Subtraction\n");
    printf("*\t: Multiplication\n");
    printf("\\\t: Division\n");
    printf("^ \t: Power\n");
}

// adds two integers
void add(int i, int j) {
    int res = i + j;
    printf("Result: %d\n", res);
}

// multiplies two numbers
int mul(int i, int j) {
    int k, res = 0;
    for (k=0; k<j; k++) {
        res += i;
    }
    printf("Result: %d\n", res);
    return res;
}

```

Attendance

Opcodes and C

When we write assembly code we are free to layout our opcodes and use registers in any way we like.

- We can place labels anywhere in our opcodes.
- We can specify a jump to any arbitrary location.
- While we can use processor support for passing return address via instructions, like `call` and `ret`, we are not required too.
- We are not forced to use the registers in any particular way.

C Standardizes how to organize and write opcodes

C forces us to decompose and organize opcodes into “functions”

- global label - single entry point
- block of opcodes ending in a “return”
- standardizes use of registers
- standardizes use of stack *use of stack allocate at the memory*
 - call frames - automatic storage of locals
 - separation into declaration (many) and definition (one)
 - compiler can get smart and optimize functions and variables away
 - in-line : *technique that eliminate jump and return*
 - dead-code elimination : *compiler can figure out any codes that never called. remove it from assembly*
 - register only variables

Let's take a look at the assembly code to see how C works

```
void myfunc(void) {}
```

returns nothing, nothing as a parameter,

```
gcc -fno-inline -fno-stack-protector -fno-pic -static -Werror -fcf-  
protection=none -fno-asynchronous-unwind-tables -Os -S -masm=intel  
myfunc0.c -o myfunc0.s
```

compile

Let's take a look at the assembly code to see how C works

```
void myfunc(void) {}
```

```
gcc -fno-inline -fno-stack-protector -fno-pic -static -Werror -fcf-  
protection=none -fno-asynchronous-unwind-tables -Os -S -masm=intel  
myfunc0.c -o myfunc0.s
```

```
.file "myfunc0.c"  
.intel_syntax noprefix  
.text  
.globl myfunc → Compiler create global Entry.  
.type myfunc, @function  
myfunc:  
ret  
.size myfunc, .-myfunc  
.ident "GCC: (Ubuntu 9.4.0-1ubuntu1~20.04.1) 9.4.0"  
.section .note.GNU-stack,"",@progbits
```

Calling a function

```
__attribute__((noinline))
```

2 functions

A int funcA(void) {
 return 7;
}

B int funcB(void) {
 return 3 + funcA();
}

```
gcc -fno-inline -fno-stack-protector -fno-pic -static -Werror -fcf-  
protection=none -fno-asynchronous-unwind-tables -Os -S -masm=intel  
myfunc1.c -o myfunc1.s
```

* EAX: return

Passing arguments

__attribute__ ((noinline))

```
int func2(int x, int y) {  
    return x + y;  
}
```

```
int func1(int x) {  
    return func2(x,2);  
}
```

```
gcc -fno-inline -fno-stack-protector -fno-pic -static -Werror -fcf-  
protection=none -fno-asynchronous-unwind-tables -Os -S -masm=intel  
myfunc3.c -o myfunc3.s
```

Passing arguments

```
__attribute__((noinline))
```

```
int func2(int x, int y) {  
    return x + y;  
}
```

```
int func1(int x) {  
    return func2(x, 2);  
}
```

```
.file "myfunc3.c"  
.intel_syntax noprefix  
.text  
.globl func2  
.type func2, @function  
func2:  
    lea eax, [rdi+rsi]  
    ret  
    .size func2, .-func2  
.globl func1  
.type func1, @function  
func1:  
    mov esi, 2  
    jmp func2  
    .size func1, .-func1  
.ident "GCC: (Ubuntu 9.4.0-1ubuntu1-20.04.1) 9.4.0"  
.section .note.GNU-stack,"",@progbits
```

load effective address
just use relocation.

let compiler to do very fast memory computation. Arithmetic operation

```
gcc -fno-inline -fno-stack-protector -fno-pic -static -Werror -fcf-  
protection=none -fno-asynchronous-unwind-tables -Os -S -masm=intel  
myfunc3.c -o myfunc3.s
```


Passing arguments

```
__attribute__((noinline))
```

```
int func2(int x, int y) {  
    return x + y;  
}
```

```
int func1(int x) {  
    return func2(x, 2);  
}
```

```
.file "myfunc3.c"  
.intel_syntax noprefix  
.text  
.globl func2  
.type func2, @function  
func2:  
    lea eax, [rdi+rsi]  
    ret  
    .size func2, .-func2  
.globl func1  
.type func1, @function  
func1:  
    mov esi, 2  
    jmp func2  
    .size func1, .-func1  
.ident "GCC: (Ubuntu 9.4.0-1ubuntu1-20.04.1) 9.4.0"  
.section .note.GNU-stack,"",@progbits
```

Why these registers?

```
gcc -fno-inline -fno-stack-protector -fno-pic -static -Werror -fcf-  
protection=none -fno-asynchronous-unwind-tables -Os -S -masm=intel  
myfunc3.c -o myfunc3.s
```

A fixed way for passing arguments

- First 6 arguments

will be
assigned to
these registers

rdi
rsi
rdx
rcx
r8
r9

more than 6th
Argument will be
located in stack

- Return value

rax

...
Arg n
...
Arg 8
Arg 7

8
7



6th rdx
rdi
rcx
rsi
r8
r9

rax

- Only allocate stack space when needed

Diane Sips Delicious Coffee 8 out of 9 times

Diane sips Delicious Coffee 8 out of 9 times.

Register saving conventions

- When `foo` calls `who`:
 - `foo` is the *caller* 호르 function A (caller is responsible for saving before to call)
 - `who` is the *callee* 호르 function B (put into stack)
- Can a register be used for temporary storage?
- Conventions
 - “Caller Saved”
 - Caller saves temporary values in its frame before the call
 - “Callee Saved”
 - Callee saves temporary values in its frame before using
 - Callee restores them before returning to caller

Stack-based Languages

we need data structure to save
callee, caller variable

"Stack"

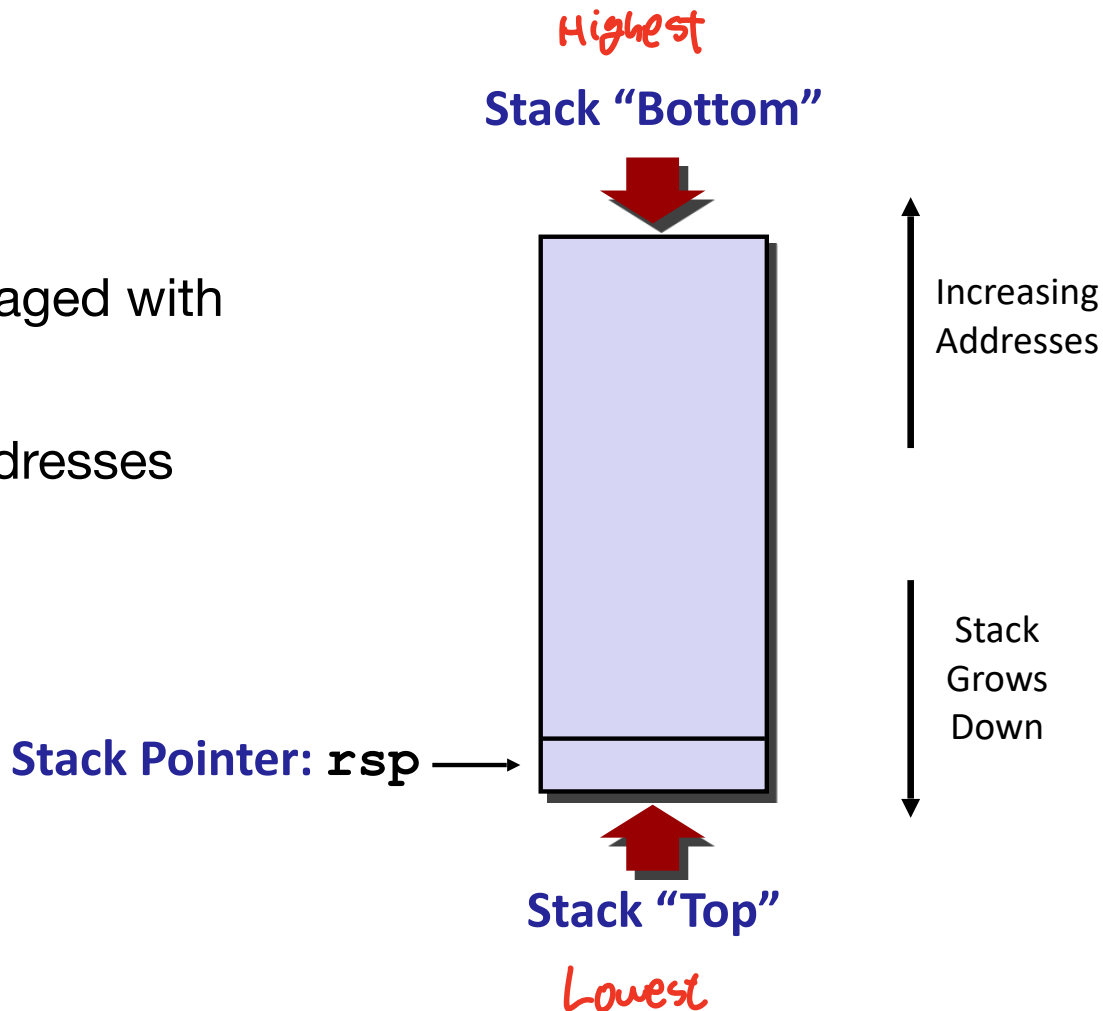
- Languages that support recursion
 - e.g., C, Pascal, Java
 - Code must be "Reentrant"
 - Multiple simultaneous instantiations of single function
 - Need some place to store state of each instantiation
 - Arguments
 - Local variables
 - Return pointer
- Stack discipline
 - State for a function is needed for limited time
 - From when called to when return
 - Callee returns before caller does
- Stack allocated in **[Frames]**
 - state for single function instantiation

F → Q

Q returns before F

x86-64 Stack


- Region of memory managed with stack discipline
- Grows toward lower addresses
- Register `rsp` contains lowest stack address
 - address of “top” element



x86-64 Stack: Push

Allocate

- push *Src*
 - Fetch operand at *Src*
 - Decrement **rsp** by 8
 - Write operand at address given by **rsp**




x86-64 Stack: Push

deallocate (return)

Stack "Bottom"

- push *Src*
 - Fetch operand at *Src*
 - Decrement **rsp** by 8
 - Write operand at address given by **rsp**

Stack Pointer: **rsp**

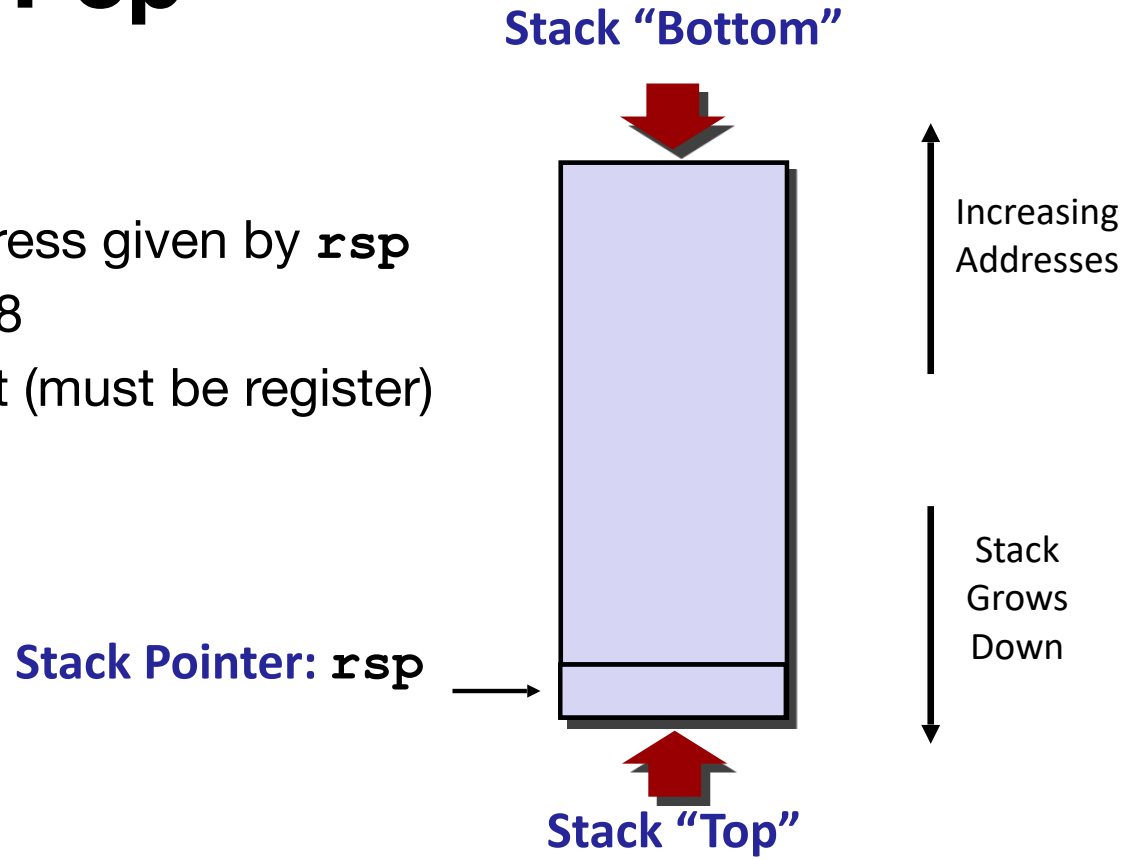


Stack "Top"



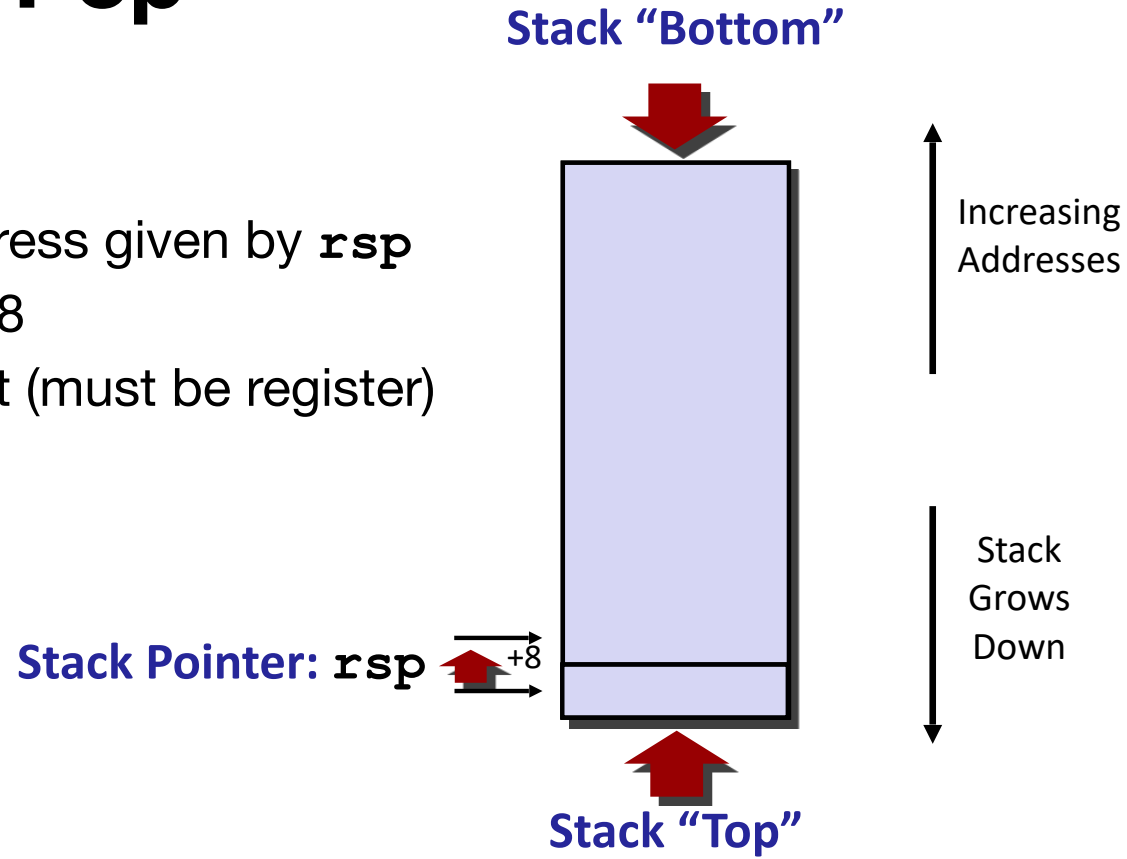
x86-64 Stack: Pop

- `pop Dest`
 - Read value at address given by `rsp`
 - Increment `rsp` by 8
 - Store value at `Dest` (must be register)



x86-64 Stack: Pop

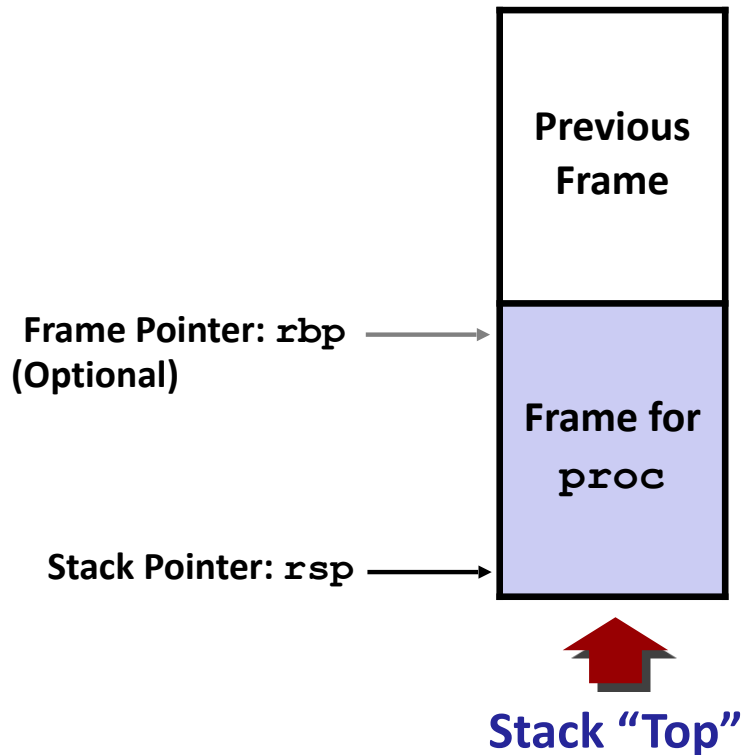
- `pop Dest`
 - Read value at address given by `rsp`
 - Increment `rsp` by 8
 - Store value at `Dest` (must be register)



Stack Frames

- Contents
 - Return information
 - Local storage (if needed)
 - Temporary space (if needed)
- Management
 - Space allocated when entering the function
 - “Set-up” code
 - Includes push by **call** instruction
 - Deallocated when return
 - “Finish” code
 - Includes pop by **ret** instruction

call includes push

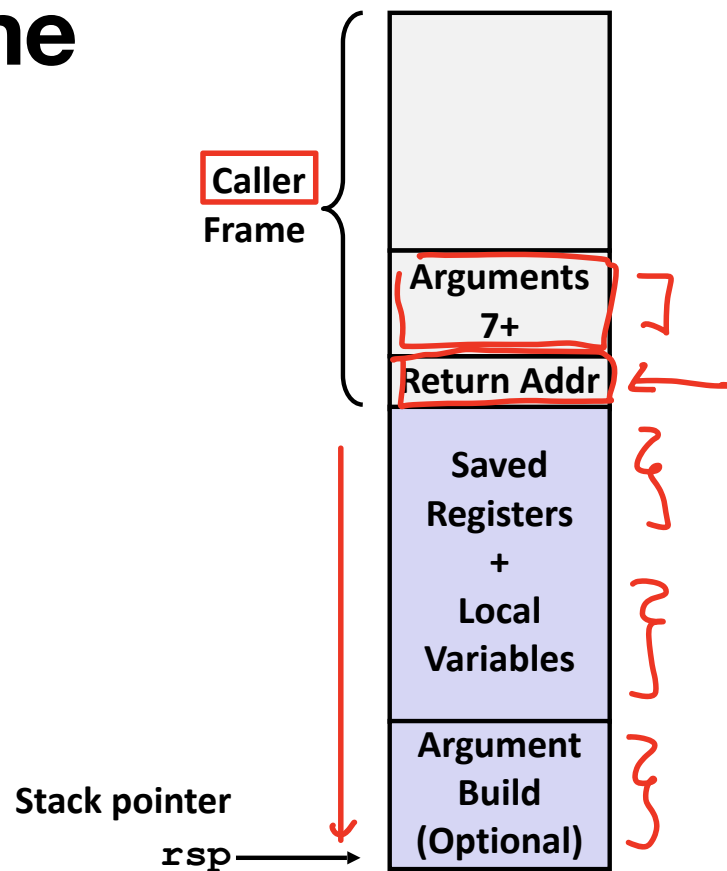


Function Control Flow

- Use stack to support call and return
- **Function call:** `call label`
 - Push return address on stack
 - Jump to ***label***
- Return address:
 - Address of the next instruction right after call
- **Function return:** `ret`
 - Pop address from stack
 - Jump to address


x86-64/Linux Stack Frame

- Current Stack Frame (“Top” to Bottom)
 - “Argument build:”
Parameters for function about to call
 - Local variables
If can’t keep in registers
 - Saved register context
- Caller Stack Frame
 - Return address
 - Pushed by `call` instruction
 - Arguments for this call




x86-64 Linux Register Usage #1

- `rax`
 - Return value
 - Also caller-saved
 - Can be modified by function
- `rdi, ..., r9`
 - Arguments
 - Also caller-saved
 - Can be modified by function
- `r10, r11`
 - Caller-saved
 - Can be modified by function



x86-64 Linux Register Usage #2

- `rbx, r12, r13, r14`
 - Callee-saved
 - Callee must save & restore
- `rbp`
 - Callee-saved
 - Callee must save & restore
 - May be used as frame pointer
 - Can mix & match
- `rsp`
 - Special form of callee save
 - Restored to original value upon exit from function




Recursion

- A function is *recursive* if it calls itself.
- The following function computes $n!$ recursively:

```
int fact(int n)
{
    if (n <= 1)
        return 1;
    else
        return n * fact(n - 1);
}
```

calling itself

Recursion



- To see how recursion works, let's trace the execution of the statement
`i = fact(3);`

`fact(3)` finds that 3 is not less than or equal to 1, so it calls
`fact(2)`, which finds that 2 is not less than or equal to 1, so it calls
`fact(1)`, which finds that 1 is less than or equal to 1, so it returns 1, causing
`fact(2)` to return $2 \times 1 = 2$, causing
`fact(3)` to return $3 \times 2 = 6$.

Call Chain Example


```
foo (...)  
{  
  .  
  .  
  who ();  
}
```

```
who (...)  
{  
  . . .  
  1 amI ();  
  . . .  
  2 amI ();  
  . . .  
}
```




```
amI (...)  
{  
  .  
  .  
  amI ();  
  .  
  .  
}
```

Procedure amI () is recursive




Example Call Chain






Example






Example






Example






Example






Example






Example






Example






Example






Example



Example



Example



Observations About Recursion

- Handled Without Special Consideration
 - Stack frames mean that each function call has private storage
 - Saved registers & local variables
 - Saved return pointer
 - Register saving conventions prevent one function call from corrupting another's data
 - Unless the C code explicitly does so (e.g., buffer overflow)
 - Stack discipline follows call / return pattern by mistake
 - If P calls Q, then Q returns before P
 - Last-In, First-Out
- Also works for mutual recursion
 - P calls Q; Q calls P

Functions Summary

- Important Points
 - Stack is the right data structure for function call / return
 - If P calls Q, then Q returns before P
- Recursion (& mutual recursion) handled by normal calling conventions
 - Can safely store values in local stack frame and in callee-saved registers
 - Put function arguments at top of stack
 - Result returned in **rax**

